

Credit Scorecard

January 1, 2022

```
[4]: # import libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import scorecardpy as sc
from string import ascii_letters
from sklearn.linear_model import LogisticRegressionCV
from sklearn.metrics import confusion_matrix
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score, confusion_matrix, roc_curve
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
%matplotlib inline
```

1 Prepare the Data

```
[67]: # create a list of variable names for the data set
varList = pd.read_excel("CC_VariablesList.XLS")
varNames = varList.loc[:, "Var_Title"].to_numpy()
varNames[43] = "MATE_EDUCATION_LEVEL"
# update the variable name which has the same with a previous variable names
print(varNames[43])
```

MATE_EDUCATION_LEVEL

```
[68]: # read in data
df = pd.read_csv("CC_Modeling_Data.
→txt", header=None, encoding="unicode_escape", delimiter="\t", names=varNames)
print(df.head())
```

	ID_CLIENT	CLERK_TYPE	PAYMENT_DAY	APPLICATION_SUBMISSION_TYPE	\
0	1	C	5		Web
1	2	C	15		Carga
2	3	C	5		Web
3	4	C	20		Web
4	5	C	10		Web

	QUANT_ADDITIONAL_CARDS	POSTAL_ADDRESS_TYPE	SEX	MARITAL_STATUS	\
0	0		1 F	6	
1	0		1 F	2	
2	0		1 F	2	
3	0		1 F	2	
4	0		1 M	2	

	QUANT_DEPENDANTS	EDUCATION_LEVEL	...	FLAG_HOME_ADDRESS_DOCUMENT	FLAG_RG	\
0	1	0	...		0	0
1	0	0	...		0	0
2	0	0	...		0	0
3	0	0	...		0	0
4	0	0	...		0	0

	FLAG_CPF	FLAG_INCOME_PROOF	PRODUCT	FLAG_ACSP_RECORD	AGE	RESIDENCIAL_ZIP_3	\
0	0	0	1	N	32	595	
1	0	0	1	N	34	230	
2	0	0	1	N	27	591	
3	0	0	1	N	61	545	
4	0	0	1	N	48	235	

	PROFESSIONAL_ZIP_3	TARGET_LABEL_BAD=1
0	595	1
1	230	1
2	591	0
3	545	0
4	235	1

[5 rows x 54 columns]

W:\Tools\Anaconda3\envs\gpu\lib\site-packages\IPython\core\interactiveshell.py:3457: DtypeWarning: Columns (51,52) have mixed types.Specify dtype option on import or set low_memory=False.
exec(code_obj, self.user_global_ns, self.user_ns)

```
[69]: # check if any variables only contain a single value which would not contribute
      ↪to our logisticregression
      for i in np.arange(len(varNames)):
          distinct = len(df.iloc[:,i].unique())
          if distinct < 2:
              print(i, " ", varNames[i])
```

```
1  CLERK_TYPE
4  QUANT_ADDITIONAL_CARDS
9  EDUCATION_LEVEL
20 FLAG_MOBILE_PHONE
44 FLAG_HOME_ADDRESS_DOCUMENT
45 FLAG_RG
```

```

46 FLAG_CPF
47 FLAG_INCOME_PROOF
49 FLAG_ACSP_RECORD

```

```

[70]: # remove those unhelpful variables and identity based variables
my_columns = np.r_[2:4,5:6,7:9,10:20,21:44,48,50:len(df.columns)]
df2 = df.iloc[:,my_columns]
print("Now there are ",len(df2.columns), " variables left")

# take a look at the distributions of the variables in general
df2.describe()

```

Now there are 43 variables left

```

[70]:      PAYMENT_DAY  POSTAL_ADDRESS_TYPE  MARITAL_STATUS  QUANT_DEPENDANTS \
count  50000.000000      50000.000000      50000.00000      50000.000000
mean    12.869920          1.006540          2.14840          0.650520
std      6.608385          0.080606          1.32285          1.193655
min       1.000000          1.000000          0.00000          0.000000
25%      10.000000          1.000000          1.00000          0.000000
50%      10.000000          1.000000          2.00000          0.000000
75%      15.000000          1.000000          2.00000          1.000000
max      25.000000          2.000000          7.00000          53.000000

```

```

      NACIONALITY  RESIDENCE_TYPE  MONTHS_IN_RESIDENCE  FLAG_EMAIL \
count  50000.000000      48651.000000      46223.000000  50000.000000
mean    0.961600          1.252225          9.727149          0.802280
std      0.202105          0.867833          10.668841          0.398284
min       0.000000          0.000000          0.000000          0.000000
25%      1.000000          1.000000          1.000000          1.000000
50%      1.000000          1.000000          6.000000          1.000000
75%      1.000000          1.000000          15.000000          1.000000
max       2.000000          5.000000          228.000000          1.000000

```

```

      PERSONAL_MONTHLY_INCOME  OTHER_INCOMES  ...  PERSONAL_ASSETS_VALUE \
count      50000.000000      50000.000000  ...      5.000000e+04
mean        886.678437        35.434760  ...      2.322372e+03
std       7846.959327       891.515142  ...      4.235798e+04
min          60.000000          0.000000  ...      0.000000e+00
25%        360.000000          0.000000  ...      0.000000e+00
50%        500.000000          0.000000  ...      0.000000e+00
75%        800.000000          0.000000  ...      0.000000e+00
max     959000.000000     194344.000000  ...      6.000000e+06

```

```

      QUANT_CARS  MONTHS_IN_THE_JOB  PROFESSION_CODE  OCCUPATION_TYPE \
count  50000.000000      50000.000000      42244.000000      42687.000000
mean    0.336140          0.009320          8.061784          2.484316
std     0.472392          0.383453          3.220104          1.532261

```

min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	9.000000	1.000000
50%	0.000000	0.000000	9.000000	2.000000
75%	1.000000	0.000000	9.000000	4.000000
max	1.000000	35.000000	18.000000	5.000000

	MATE_PROFESSION_CODE	MATE_EDUCATION_LEVEL	PRODUCT	AGE \
count	21116.000000	17662.000000	50000.000000	50000.00000
mean	3.797926	0.296003	1.275700	43.24852
std	5.212168	0.955688	0.988286	14.98905
min	0.000000	0.000000	1.000000	6.00000
25%	0.000000	0.000000	1.000000	31.00000
50%	0.000000	0.000000	1.000000	41.00000
75%	11.000000	0.000000	1.000000	53.00000
max	17.000000	5.000000	7.000000	106.00000

	TARGET_LABEL_BAD=1
count	50000.000000
mean	0.260820
std	0.439086
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

[8 rows x 27 columns]

2 Data Cleaning

2.1 check for Null values

```
[71]: # find out which columns have null values
df2.isnull().any()
```

```
[71]: PAYMENT_DAY                False
APPLICATION_SUBMISSION_TYPE     False
POSTAL_ADDRESS_TYPE             False
MARITAL_STATUS                  False
QUANT_DEPENDANTS                False
STATE_OF_BIRTH                  False
CITY_OF_BIRTH                   False
NACIONALITY                     False
RESIDENCIAL_STATE               False
RESIDENCIAL_CITY                False
RESIDENCIAL_BOROUGH             False
FLAG_RESIDENCIAL_PHONE          False
```

RESIDENCIAL_PHONE_AREA_CODE	False
RESIDENCE_TYPE	True
MONTHS_IN_RESIDENCE	True
FLAG_EMAIL	False
PERSONAL_MONTHLY_INCOME	False
OTHER_INCOMES	False
FLAG_VISA	False
FLAG_MASTERCARD	False
FLAG_DINERS	False
FLAG_AMERICAN_EXPRESS	False
FLAG_OTHER_CARDS	False
QUANT_BANKING_ACCOUNTS	False
QUANT_SPECIAL_BANKING_ACCOUNTS	False
PERSONAL_ASSETS_VALUE	False
QUANT_CARS	False
COMPANY	False
PROFESSIONAL_STATE	False
PROFESSIONAL_CITY	True
PROFESSIONAL_BOROUGH	True
FLAG_PROFESSIONAL_PHONE	False
PROFESSIONAL_PHONE_AREA_CODE	False
MONTHS_IN_THE_JOB	False
PROFESSION_CODE	True
OCCUPATION_TYPE	True
MATE_PROFESSION_CODE	True
MATE_EDUCATION_LEVEL	True
PRODUCT	False
AGE	False
RESIDENCIAL_ZIP_3	False
PROFESSIONAL_ZIP_3	False
TARGET_LABEL_BAD=1	False

dtype: bool

```
[72]: # find how many null values those columns have
null_columns = df2.columns[df2.isnull().any()]
df2[null_columns].isnull().sum()
```

```
[72]: RESIDENCE_TYPE          1349
MONTHS_IN_RESIDENCE        3777
PROFESSIONAL_CITY          33783
PROFESSIONAL_BOROUGH       33783
PROFESSION_CODE            7756
OCCUPATION_TYPE            7313
MATE_PROFESSION_CODE       28884
MATE_EDUCATION_LEVEL       32338
dtype: int64
```

```
[73]: # count how many " " and Null values in each row, and count how many " " in
      ↪ each column
null_rows=np.zeros(len(df2))
missing_col=np.zeros(len(df2.columns))
for k in np.arange(len(df2)):
    count=0
    for i in np.arange(len(df2.columns)):
        val = df2.iloc[k,i]
        #print(val)
        if val == " ":
            missing_col[i] += 1
            count += 1
    count += df2.loc[k,:].isnull().sum()
    null_rows[k]=count
```

```
[74]: # find out if any row has more than 30% null values (13 null values)

null_rows2 = null_rows/len(df2.columns) > 0.3
print(df2.iloc[null_rows2,-1])
# I shouldn't delete any rows
```

Series([], Name: TARGET_LABEL_BAD=1, dtype: int64)

```
[75]: # combine the results for missing values
comb_miss = np.zeros(len(missing_col))

for i in np.arange(len(missing_col)):
    comb_miss[i] = missing_col[i]
    for j in np.arange(len(null_columns)):
        if df2.columns[i] == null_columns[j]:
            comb_miss[i] += df2[null_columns[j]].isnull().sum()
        continue
# print the final result of missing values statistics
for i in np.arange(len(comb_miss)):
    if comb_miss[i] != 0:
        print("{0:<40s}{1}".format(df2.columns[i],int(comb_miss[i])))
```

STATE_OF_BIRTH	2064
CITY_OF_BIRTH	2064
RESIDENCIAL_BOROUGH	10
RESIDENCIAL_PHONE_AREA_CODE	8212
RESIDENCE_TYPE	1349
MONTHS_IN_RESIDENCE	3777
PROFESSIONAL_STATE	34307
PROFESSIONAL_CITY	34114
PROFESSIONAL_BOROUGH	34713
PROFESSIONAL_PHONE_AREA_CODE	36532
PROFESSION_CODE	7756

OCCUPATION_TYPE	7313
MATE_PROFESSION_CODE	28884
MATE_EDUCATION_LEVEL	32338

2.2 missing value treatment #1

```
[76]: # variables to delete:
# professional state, professional city, professional borough, and
# professional phone area code
# which have more than 60% missing values, the information is not important
# also delete mate profession code and mate education level because more than
# 50% missing value and information
# not important
# delete residential phone area code, has 17% missing value and use falg
# residencial phone as dummy variable for it

df2=df2.drop(axis=1,
# columns=["PROFESSIONAL_STATE","PROFESSIONAL_CITY","PROFESSIONAL_BOROUGH","PROFESSIONAL_PHONE_AREA_CODE",
# "MATE_PROFESSION_CODE","MATE_EDUCATION_LEVEL","RESIDENCIAL_PHONE_AREA_CODE"])
```

```
[77]: # replace recordings of residence type 0 as 1
df2["RESIDENCE_TYPE"].replace({0:1},inplace=True)
# now,create a category for valid null values in residence type
df2.RESIDENCE_TYPE.fillna(value = 0,inplace=True)

# now,create a category for valid null values in profession code and occupation
# type
df2.PROFESSION_CODE.fillna(value = 19,inplace=True)
df2.OCCUPATION_TYPE.fillna(value = 6,inplace=True)
```

```
[78]: print("There are ", len(df2.columns), " variables left")
```

There are 36 variables left

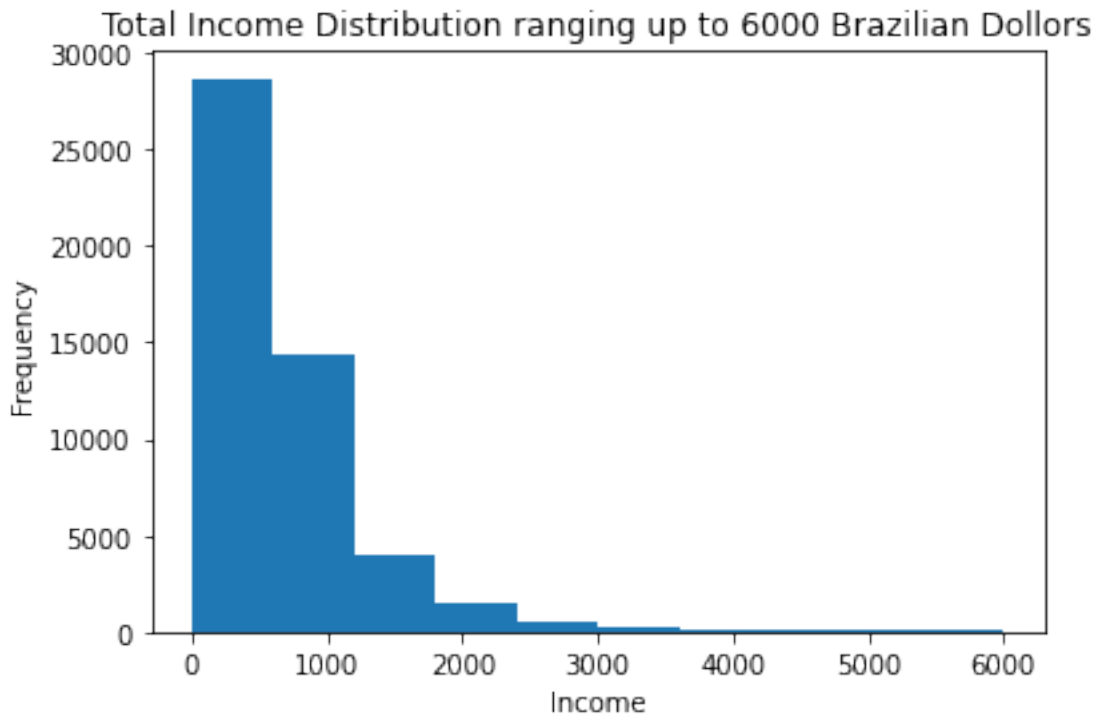
2.3 treating outliers in total income


```
[79]: df2["TOTAL_INCOME"] = df2["PERSONAL_MONTHLY_INCOME"]+df2["OTHER_INCOMES"]
df2.TOTAL_INCOME.describe()
```

```
[79]: count    50000.000000
mean       922.113196
std        7897.469079
min         60.000000
25%        372.000000
50%        515.000000
75%        840.000000
```

```
max      959000.000000  
Name: TOTAL_INCOME, dtype: float64
```

```
[80]: # take a look at income  
plt.hist(df2["TOTAL_INCOME"],range=(0,6000))  
#plt.savefig("Income.pdf")  
plt.xlabel("Income")  
plt.ylabel("Frequency")  
plt.title("Total Income Distribution ranging up to 6000 Brazilian Dollors")  
plt.show()  
plt.hist(df2["TOTAL_INCOME"],range=(0,4000))  
#plt.savefig("Income2.pdf")  
plt.xlabel("Income")  
plt.ylabel("Frequency")  
plt.title("Total Income Distribution ranging up to 4000 Brazilian Dollors")  
plt.show()
```





```
[81]: # only consider customers with total income less than or equal to 4000
df2 = df2.loc[df2["TOTAL_INCOME"] <= 4000]
```

2.4 make new variables

```
[82]: # variable 1: supporting power of dependents and self
df2["SUPPORT_POWER"] = df2["TOTAL_INCOME"] / (df2["QUANT_DEPENDANTS"] + 1)

# variable 2: accomplishment
df2["ACCOMPLISHMENT"] = df2["TOTAL_INCOME"] / df2["AGE"]

# variable 3: debt over standardized income
mean_income = df2.TOTAL_INCOME.mean()
std_income = df2.TOTAL_INCOME.std()
df2["DEBT"] =
    ↪ (df2["FLAG_VISA"] + df2["FLAG_MASTERCARD"] + df2["FLAG_DINERS"] + df2["FLAG_AMERICAN_EXPRESS"] + df2["FLAG_OTHER_CREDIT_CARDS"]) / (std_income * 1000)
```

2.5 split the data and replace missing values

```
[83]: # split the data into training and testing set with 70-30 ratio

train, test = sc.split_df(df2, y="TARGET_LABEL_BAD=1", ratio=0.7, seed=250918939).
    ↪ values()
```

W:\Tools\Anaconda3\envs\gpu\lib\site-packages\scorecardpy\condition_fun.py:62:
 UserWarning: There are blank strings in 3 columns, which are replaced with NaN.
 (ColumnNames: STATE_OF_BIRTH, CITY_OF_BIRTH, RESIDENCIAL_BOROUGH)
 warnings.warn('There are blank strings in {} columns, which are replaced with
 NaN. \n (ColumnNames: {})).format(len(blank_cols), ', '.join(blank_cols)))

```
[84]: # replace state of birth and city of birth by mode from the training set
      ↪(replace missing values instead of null values)
temp_col = train.STATE_OF_BIRTH[train.STATE_OF_BIRTH != " "]
birthstate_mode = np.array(temp_col.mode())[0]
train["STATE_OF_BIRTH"].replace({" ":birthstate_mode},inplace=True)
test["STATE_OF_BIRTH"].replace({" ":birthstate_mode},inplace=True)
df2["STATE_OF_BIRTH"].replace({" ":birthstate_mode},inplace=True)

temp_col = train.CITY_OF_BIRTH[train.CITY_OF_BIRTH != " "]
birthcity_mode = np.array(temp_col.mode())[0]
train["CITY_OF_BIRTH"].replace({" ":birthcity_mode},inplace=True)
test["CITY_OF_BIRTH"].replace({" ":birthcity_mode},inplace=True)
df2["CITY_OF_BIRTH"].replace({" ":birthcity_mode},inplace=True)

# replace residential borough by the mode from the traing set (missing values)
temp_col = train.RESIDENCIAL_BOROUGH[train.RESIDENCIAL_BOROUGH != " "]
residencialb_mode = np.array(temp_col.mode())[0]
train["RESIDENCIAL_BOROUGH"].replace({" ":residencialb_mode},inplace=True)
test["RESIDENCIAL_BOROUGH"].replace({" ":residencialb_mode},inplace=True)
df2["RESIDENCIAL_BOROUGH"].replace({" ":residencialb_mode},inplace=True)

# replace month in residence by the median (null values)
train.MONTHS_IN_RESIDENCE.fillna(value=train.MONTHS_IN_RESIDENCE.
    ↪median(),inplace=True)
test.MONTHS_IN_RESIDENCE.fillna(value=train.MONTHS_IN_RESIDENCE.
    ↪median(),inplace=True)
df2.MONTHS_IN_RESIDENCE.fillna(value=train.MONTHS_IN_RESIDENCE.
    ↪median(),inplace=True)

[85]: # replace age < 18 by the median since those people won't be able to get a loan
age_median=train.AGE.median()
df2["AGE"] = df2.apply(lambda x: age_median if x["AGE"] < 18 else x["AGE"],
    ↪axis=1)
train["AGE"] = train.apply(lambda x: age_median if x["AGE"] < 18 else x["AGE"],
    ↪axis=1)
test["AGE"] = test.apply(lambda x: age_median if x["AGE"] < 18 else x["AGE"],
    ↪axis=1)

[86]: # check if null values in the training set have been delt with
missing_col=np.zeros(len(df2.columns))
```

```

for k in np.arange(len(train)):
    count=0
    for i in np.arange(len(df2.columns)):
        val = train.iloc[k,i]
        #print(val)
        if val == " ":
            missing_col[i] += 1

# check the columns with " " value
for i in np.arange(len(missing_col)):
    if missing_col[i] != 0:
        print("{0:<30s}{1}".format(train.columns[i],int(missing_col[i])))

# recheck the columns with null values
null_columns = train.columns[train.isnull().any()]
train[null_columns].isnull().sum()

```

[86]: Series([], dtype: float64)

3 Weight of Evidence Transformation and Logistic Model

3.1 combining small categories

```

[87]: # standardizing cities nad boroughs
df2["CITY_OF_BIRTH"] = df2.apply(lambda x: (x["CITY_OF_BIRTH"].replace(" ", "").upper(), axis=1)
df2["RESIDENCIAL_CITY"] = df2.apply(lambda x: (x["RESIDENCIAL_CITY"].replace(" ", "").upper(), axis=1)
df2["RESIDENCIAL_BOROUGH"] = df2.apply(lambda x: (x["RESIDENCIAL_BOROUGH"].replace(" ", "").upper(), axis=1)

```

```

[88]: # stadnardizing the codes
df2["RESIDENCIAL_ZIP_3"] = df2.apply(lambda x: int(x["RESIDENCIAL_ZIP_3"]) if (
    (type(x["RESIDENCIAL_ZIP_3"])==int or type(x["RESIDENCIAL_ZIP_3"])==str)
    and x["RESIDENCIAL_ZIP_3"] != "#DIV/0!") else 0, axis=1 )
df2["PROFESSIONAL_ZIP_3"] = df2.apply(lambda x: int(x["PROFESSIONAL_ZIP_3"]) if (
    (type(x["PROFESSIONAL_ZIP_3"])==int or type(x["PROFESSIONAL_ZIP_3"])==str)
    and x["PROFESSIONAL_ZIP_3"] != "#DIV/0!") else 0, axis=1 )

```

```

[89]: # define a function that takes dataframe, varaible and return the categories
      #with small counts
def retrieve(mincase,df,var):
    values = df[var].value_counts()
    values2 = pd.DataFrame(data=values)
    values2 = values2[values2[var]<mincase]
    values3 = np.array(values2.index)

```

```
return values3
```

```
[91]: # find categories with small amount of cases and combine them
mincase=50
smallbirth_cities = retrieve(mincase,df2,"CITY_OF_BIRTH")
#print(smallbirth_cities)

smallres_cities = retrieve(mincase, df2, "RESIDENCIAL_CITY")

smallres_borough = retrieve(mincase, df2, "RESIDENCIAL_BOROUGH")

smallres_zip = retrieve(mincase, df2, "RESIDENCIAL_ZIP_3")

smallpro_zip = retrieve(mincase, df2, "PROFESSIONAL_ZIP_3")

[92]: # make sure that each category in categorical variables have good amount of
      ↪ cases
df2["CITY_OF_BIRTH"] = df2.apply(lambda x: "OTHERS" if x["CITY_OF_BIRTH"] in
      ↪ smallbirth_cities else x["CITY_OF_BIRTH"],
                                axis=1)
df2["RESIDENCIAL_CITY"] = df2.apply(lambda x: "OTHERS" if x["RESIDENCIAL_CITY"]
      ↪ in smallres_cities else x["RESIDENCIAL_CITY"],
                                axis=1)
df2["RESIDENCIAL_BOROUGH"] = df2.apply(lambda x: "OTHERS" if
      ↪ x["RESIDENCIAL_BOROUGH"] in smallres_borough else x["RESIDENCIAL_BOROUGH"],
                                axis=1)
df2["RESIDENCIAL_ZIP_3"] = df2.apply(lambda x: 0 if x["RESIDENCIAL_ZIP_3"] in
      ↪ smallres_zip else x["RESIDENCIAL_ZIP_3"],
                                axis=1)
df2["PROFESSIONAL_ZIP_3"] = df2.apply(lambda x: 0 if x["PROFESSIONAL_ZIP_3"] in
      ↪ smallpro_zip else x["PROFESSIONAL_ZIP_3"],
                                axis=1)
```

3.2 create bins

```
[94]: # split again
train, test = sc.split_df(df2,y="TARGET_LABEL_BAD=1",ratio=0.7,seed=250918939).
      ↪ values()

[96]: # with some testing, decided to drop the variables with 0 and almost 0
      ↪ information values
df3 = df2.loc[:,["OCCUPATION_TYPE","RESIDENCIAL_STATE","MARITAL_STATUS",
                "RESIDENCIAL_ZIP_3",
                "AGE", "PAYMENT_DAY",
      ↪ "PROFESSIONAL_ZIP_3","ACCOMPLISHMENT","TARGET_LABEL_BAD=1"]]
train = train.loc[:,["OCCUPATION_TYPE","RESIDENCIAL_STATE","MARITAL_STATUS",
                    "RESIDENCIAL_ZIP_3",
```

```

        "AGE", "PAYMENT_DAY", "PROFESSIONAL_ZIP_3", "ACCOMPLISHMENT",
        "TARGET_LABEL_BAD=1"]]]
test = test.loc[:, ["OCCUPATION_TYPE", "RESIDENCIAL_STATE", "MARITAL_STATUS",
        "RESIDENCIAL_ZIP_3",
        "AGE", "PAYMENT_DAY", "PROFESSIONAL_ZIP_3", "ACCOMPLISHMENT",
        "TARGET_LABEL_BAD=1"]]

```

```

[97]: # Try with 100 cuts, minimum 2% data(1000 cases) per bin and maximum 7 bins
bins3 = sc.woebin(train, y="TARGET_LABEL_BAD=1",
        min_perc_fine_bin=0.01,
        min_perc_coarse_bin=0.05,
        stop_limit=0.02,
        max_num_bin=7,
        method="tree"
        )

```

[INFO] creating woe binning ...

```

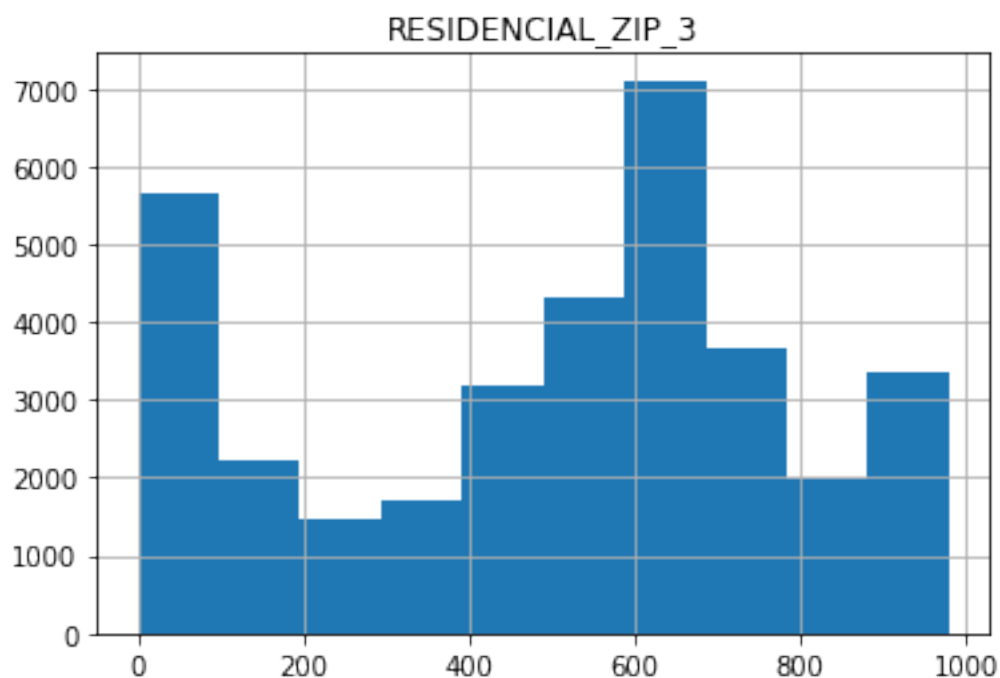
[99]: # adjust bins so that they represent meaningful trends
breaks_adj = sc.woebin_adj(train, "TARGET_LABEL_BAD=1", bins3, adj_all_var=True)

```


```

----- 1/8 RESIDENCIAL_ZIP_3 -----
>>> dt[RESIDENCIAL_ZIP_3].describe():
count      34680.000000
mean         493.526413
std          295.999862
min           0.000000
25%          286.000000
50%          572.000000
75%          689.000000
max          980.000000
Name: RESIDENCIAL_ZIP_3, dtype: float64

```



```
>>> Current breaks:
130.0,400.0,610.0,690.0,750.0,950.0
```



```
>>> Adjust breaks for (1/8) RESIDENCIAL_ZIP_3?
```

```
1: next
```

```
2: yes
```

```
3: back
```

```
Selection: 1
```

```
----- 2/8 MARITAL_STATUS -----
```

```
>>> dt[MARITAL_STATUS].describe():
```

```
count      34680.000000
```

```
mean         2.143829
```

```
std          1.317922
```

```
min          0.000000
```

```
25%          1.000000
```

```
50%          2.000000
```

```
75%          2.000000
```

```
max          7.000000
```

```
Name: MARITAL_STATUS, dtype: float64
```

```
>>> dt[MARITAL_STATUS].value_counts():
```

```
2      17986
```

```
1      10678
```

```
4       2927
```

```
6       1313
```

```
5        871
```

```
3        431
```


```
7        352
```

```
0        122
```

```
Name: MARITAL_STATUS, dtype: int64
```

```
>>> Current breaks:
```

```
2.0,3.0,5.0
```



```
>>> Adjust breaks for (2/8) MARITAL_STATUS?
```

```
1: next
```

```
2: yes
```

```
3: back
```

```
Selection: 1
```

```
----- 3/8 PROFESSIONAL_ZIP_3 -----
```

```
>>> dt[PROFESSIONAL_ZIP_3].describe():
```

```
count    34680.000000
```

```
mean      493.526413
```

```
std       295.999862
```

```
min        0.000000
```

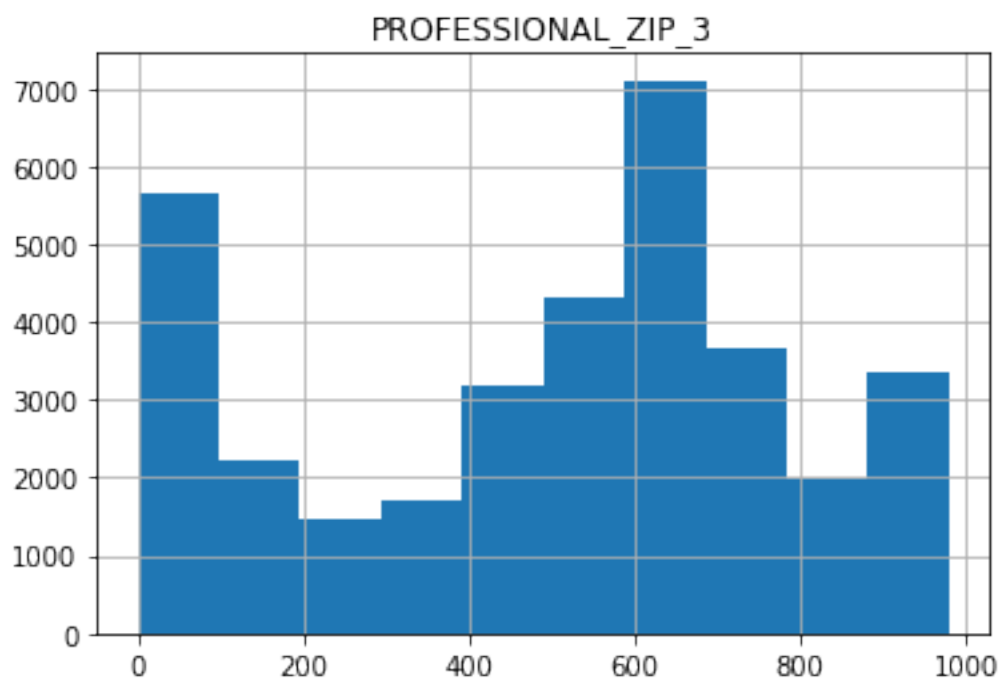
```
25%       286.000000
```

```
50%       572.000000
```


```
75%       689.000000
```

```
max       980.000000
```

```
Name: PROFESSIONAL_ZIP_3, dtype: float64
```

```
>>> Current breaks:
130.0,400.0,610.0,690.0,750.0,950.0
```



```
>>> Adjust breaks for (3/8) PROFESSIONAL_ZIP_3?
```

```
1: next
```

```
2: yes
```

```
3: back
```

```
Selection: 1
```

```
----- 4/8 PAYMENT_DAY -----
```

```
>>> dt[PAYMENT_DAY].describe():
```

```
count      34680.000000
```

```
mean        12.889937
```

```
std         6.602135
```

```
min         1.000000
```

```
25%        10.000000
```

```
50%        10.000000
```

```
75%        15.000000
```

```
max        25.000000
```

```
Name: PAYMENT_DAY, dtype: float64
```

```
>>> dt[PAYMENT_DAY].value_counts():
```

```
10      14109
```

```
15       6110
```

```
5        5106
```

```
25       4864
```


```
20       3298
```

```
1        1193
```


```
Name: PAYMENT_DAY, dtype: int64
```

```
>>> Current breaks:
```

```
10.0,15.0,20.0,25.0
```



```
>>> Adjust breaks for (4/8) PAYMENT_DAY?
1: next
2: yes
3: back
Selection: 2
>>> Enter modified breaks: 11,18
[INFO] creating woe binning ...
>>> Current breaks:
11.0, 18.0
```



```
>>> Adjust breaks for (4/8) PAYMENT_DAY?
```

```
1: next
```

```
2: yes
```

```
3: back
```

```
Selection: 1
```

```
----- 5/8 AGE -----
```

```
>>> dt[AGE].describe():
```

```
count    34680.000000
```

```
mean      43.140398
```

```
std       15.050384
```

```
min       18.000000
```


```
25%       31.000000
```

```
50%       41.000000
```


```
75%       53.000000
```

```
max       106.000000
```


```
Name: AGE, dtype: float64
```



```
>>> Current breaks:
22.0,33.0,44.0,47.0,55.0,70.0
```



```
>>> Adjust breaks for (5/8) AGE?
1: next
2: yes
3: back
Selection: 2
>>> Enter modified breaks: 26,33,43,49,60
[INFO] creating woe binning ...
>>> Current breaks:
60.0, 49.0, 26.0, 33.0, 43.0
```



```
>>> Adjust breaks for (5/8) AGE?
1: next
2: yes
3: back
Selection: 1
----- 6/8 RESIDENCIAL_STATE -----
>>> dt[RESIDENCIAL_STATE].describe():
count      34680
unique       27
top         SP
freq       6127
Name: RESIDENCIAL_STATE, dtype: object
```


```
>>> dt[RESIDENCIAL_STATE].value_counts():
```

SP	6127
RS	3652
BA	3548
CE	3235
PE	2517
MG	2153
PA	1551
RJ	1469
RN	1404
GO	1131
PR	1069
AL	1001
PB	912
MT	893
MA	515
MS	513
ES	491
DF	469
SC	442
AP	335
AM	281
PI	254
RO	244
SE	184
AC	160
TO	87
RR	43

```
Name: RESIDENCIAL_STATE, dtype: int64
```

```
>>> Current breaks:
```

```
'SC%', 'RO%', 'RS', 'PR%', 'PB', 'PA%', 'MG%', 'AP', 'AC%', 'MA%', 'SP', 'MS%', 'TO%', 'RJ%', 'MT%', 'CE%',  
'PI', 'GO%', 'BA%', 'RR%', 'RN%', 'PE%', 'ES', 'AL%', 'AM%', 'DF%', 'SE'
```



```
>>> Adjust breaks for (6/8) RESIDENCIAL_STATE?
```

```
1: next
```

```
2: yes
```

```
3: back
```

```
Selection: 1
```

```
----- 7/8 OCCUPATION_TYPE -----
```

```
>>> dt[OCCUPATION_TYPE].describe():
```

```
count    34680.000000
```

```
mean      3.001557
```

```
std       1.882931
```

```
min       0.000000
```

```
25%       2.000000
```

```
50%       2.000000
```

```
75%       5.000000
```

```
max       6.000000
```

```
Name: OCCUPATION_TYPE, dtype: float64
```

```
>>> dt[OCCUPATION_TYPE].value_counts():
```

```
2.0    11865
```

```
1.0     6000
```

```
6.0     5117
```

```
4.0     4827
```

```
5.0     4743
```


```
0.0     1915
```

```
3.0      213
```


Name: OCCUPATION_TYPE, dtype: int64

>>> Current breaks:

1.0,2.0,4.0,5.0,6.0



>>> Adjust breaks for (7/8) OCCUPATION_TYPE?

1: next

2: yes

3: back

Selection: 1

----- 8/8 ACCOMPLISHMENT -----

>>> dt[ACCOMPLISHMENT].describe():

count 34680.000000

mean 17.901907

std 13.347868

min 1.754386

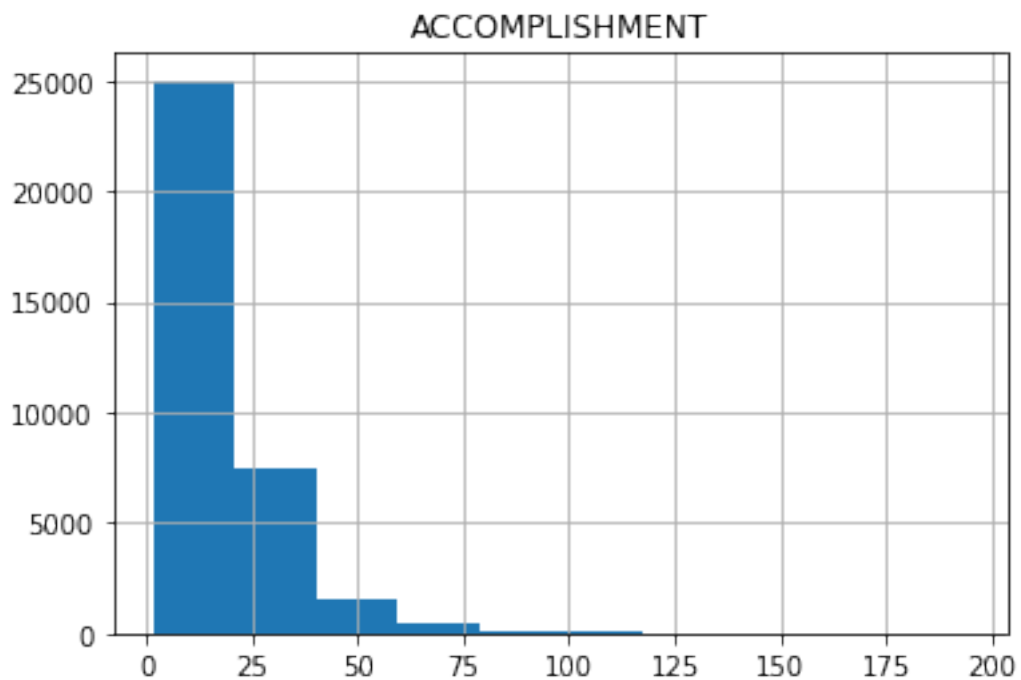
25% 9.090909

50% 14.317627


75% 22.254798

max 194.000000


Name: ACCOMPLISHMENT, dtype: float64



```
>>> Current breaks:
8.0,12.0,23.0,43.0
```



```
>>> Adjust breaks for (8/8) ACCOMPLISHMENT?
1: next
2: yes
3: back
Selection: 2
>>> Enter modified breaks: 8,12,23
[INFO] creating woe binning ...
>>> Current breaks:
8.0, 23.0, 12.0
```



```
>>> Adjust breaks for (8/8) ACCOMPLISHMENT?
1: next
2: yes
3: back
Selection: 1
```

```
[100]: # apply the adjusted cuts
bins_final = sc.woebin(train, y="TARGET_LABEL_BAD=1",breaks_list=breaks_adj)
# calculate WoE dataset
train_woe = sc.woebin_ply(train, bins_final)
test_woe = sc.woebin_ply(test, bins_final)
```

```
[INFO] creating woe binning ...
```

```
[INFO] converting into woe values ...
[INFO] converting into woe values ...
```

```
[101]: # information value
sc.iv(train_woe,"TARGET_LABEL_BAD=1")
```

```
[101]:
```

	variable	info_value
7	AGE_woe	0.067863
3	PAYMENT_DAY_woe	0.030390
4	MARITAL_STATUS_woe	0.027959
6	ACCOMPLISHMENT_woe	0.021630
2	OCCUPATION_TYPE_woe	0.018840
1	RESIDENCIAL_STATE_woe	0.016862
0	PROFESSIONAL_ZIP_3_woe	0.015906
5	RESIDENCIAL_ZIP_3_woe	0.015906

3.3 correlation check

```
[102]: # compute the correlation matrix
corr = train_woe.corr()
corr = np.abs(corr)


# generate a mask for the upper triangle correlation matrix
mask = np.triu(np.ones_like(corr,dtype=bool))

# set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# generate a diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# draw the heatmap with the mask nad correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=1, center=0, square=True,
            ↳linewidths=.5, cbar_kws={"shrink": .5})
```

```
[102]: <AxesSubplot:>
```



```
[103]: # removing residencial zip code
df3 = df2.loc[:,["OCCUPATION_TYPE","RESIDENCIAL_STATE","MARITAL_STATUS",
                "AGE", "PAYMENT_DAY",
                ↪"PROFESSIONAL_ZIP_3","ACCOMPLISHMENT","TARGET_LABEL_BAD=1"]]
train = train.loc[:,["OCCUPATION_TYPE","RESIDENCIAL_STATE","MARITAL_STATUS",
                    "AGE", "PAYMENT_DAY", "PROFESSIONAL_ZIP_3","ACCOMPLISHMENT",
                    "TARGET_LABEL_BAD=1"]]
test = test.loc[:,["OCCUPATION_TYPE","RESIDENCIAL_STATE","MARITAL_STATUS",
                  "AGE", "PAYMENT_DAY", "PROFESSIONAL_ZIP_3","ACCOMPLISHMENT",
                  "TARGET_LABEL_BAD=1"]]
# calculate WoE dataset
train_woe = sc.woebin_ply(train, bins_final)
test_woe = sc.woebin_ply(test, bins_final)
train_woe.head()
```

[INFO] converting into woe values ...

[INFO] converting into woe values ...

```
[103]: TARGET_LABEL_BAD=1  MARITAL_STATUS_woe  PROFESSIONAL_ZIP_3_woe  \
1          1          -0.094251          -0.065840
2          0          -0.094251          0.116334
3          0          -0.094251          0.116334
4          1          -0.094251          -0.065840
5          1          -0.094251          -0.065840

      PAYMENT_DAY_woe  AGE_woe  RESIDENCIAL_STATE_woe  OCCUPATION_TYPE_woe  \
1      0.059953  0.044956          0.038654          0.133313
2     -0.139198  0.272041          0.093864          0.010420
3      0.273643 -0.396786          0.093864          0.010420
4     -0.139198 -0.066754          0.038654          0.034756
5     -0.139198  0.044956         -0.059674          0.060699

      ACCOMPLISHMENT_woe
1      0.049615
2      0.049615
3     -0.098086
4      0.182665
5      0.049615
```

3.4 train the model

```
[104]: score_logreg = LogisticRegressionCV(penalty="elasticnet",
      Cs = [i for i in range(10,50)], # try 40 parameters
      tol = 0.0001,
      cv = 3,
      fit_intercept=True,
      class_weight = "balanced",
      random_state = 250918939,
      max_iter = 100,
      verbose = 0,
      solver = "saga",
      n_jobs=2,
      refit = True,
      l1_ratios = np.arange(0, 1.01, 0.1))

[105]: score_logreg.fit(X = train_woe.iloc[:,1:], y=train_woe["TARGET_LABEL_BAD=1"])

[105]: LogisticRegressionCV(Cs=[10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,
      24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37,
      38, 39, ...],
      class_weight='balanced', cv=3,
      l1_ratios=array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7,
      0.8, 0.9, 1. ]),
      n_jobs=2, penalty='elasticnet', random_state=250918939,
```

```
solver='saga')
```

```
[106]: # parameters: should not expect negative coefficients
coef_df = pd.concat([pd.DataFrame({"columns":train_woe.columns[1:]}),
                    pd.DataFrame(np.transpose(score_logreg.coef_))],
                    axis=1)
coef_df
```

```
[106]:
```

	columns	0
0	MARITAL_STATUS_woe	0.514044
1	PROFESSIONAL_ZIP_3_woe	0.590490
2	PAYMENT_DAY_woe	0.750327
3	AGE_woe	0.765557
4	RESIDENCIAL_STATE_woe	0.576933
5	OCCUPATION_TYPE_woe	-0.054264
6	ACCOMPLISHMENT_woe	0.216293

3.5 test performance

```
[108]: # apply the model to test set
pred_test = score_logreg.predict(test_woe.
    ↳drop(axis=1,columns="TARGET_LABEL_BAD=1"))
```

```
[109]: # Calculate confusion matrix
confusion_matrix_logreg =
    ↳confusion_matrix(y_true=test["TARGET_LABEL_BAD=1"],y_pred = pred_test)

# turn matrix to percentage
confusion_matrix_logreg = confusion_matrix_logreg.astype("float") /
    ↳confusion_matrix_logreg.sum(axis=1)[:,np.newaxis]

# turn to dataframe
df_cm = pd.DataFrame(confusion_matrix_logreg, index=["good", "bad"],
    ↳columns=["good", "bad"])


# parameters of the image
figsize = (10, 7)
fontsize = 14

# Crate image
fig = plt.figure(figsize=figsize)
heatmap = sns.heatmap(df_cm,annot=True,fmt=".2f")

heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0,
    ↳ha="right", fontsize=fontsize)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45,
    ↳ha="right", fontsize=fontsize)
```

```
plt.ylabel("True label")
plt.xlabel("Predicted label")

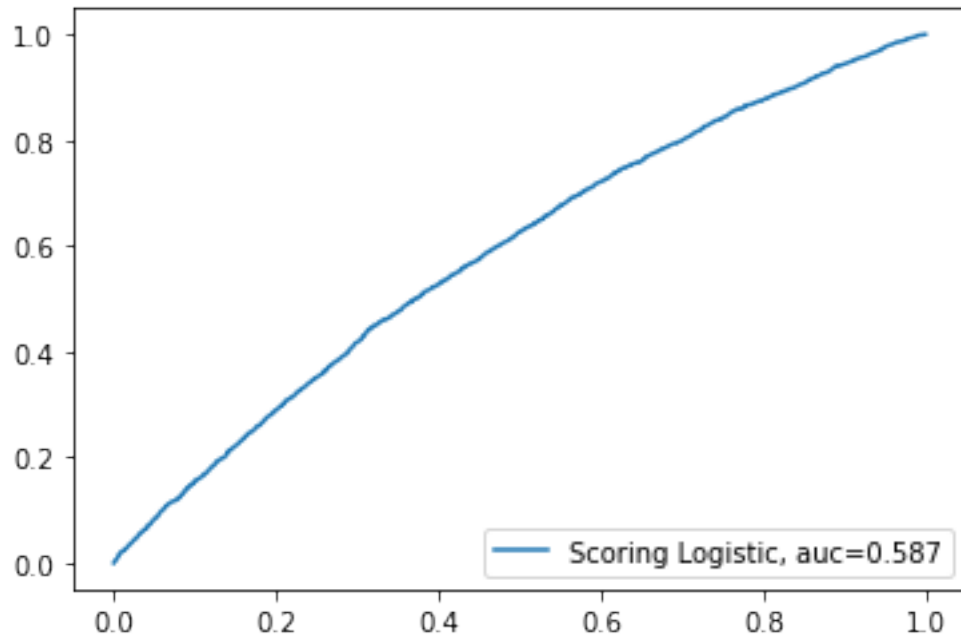
plt.show()
```



```
[110]: # calculate the ROC curve points
pred_test_prob = score_logreg.predict_proba(test_woe.
    ↳ drop(axis=1, columns="TARGET_LABEL_BAD=1"))
fpr, tpr, thresholds = roc_curve(test["TARGET_LABEL_BAD=1"], pred_test_prob[:, 1])

auc = np.round(roc_auc_score(y_true=test["TARGET_LABEL_BAD=1"],
    ↳ y_score=pred_test_prob[:, 1]), decimals=3)

plt.plot(fpr, tpr, label="Scoring Logistic, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```

```
[111]: # variable importance
sc.iv(train_woe,"TARGET_LABEL_BAD=1")
```

```
[111]:
```

	variable	info_value
6	AGE_woe	0.067863
3	PAYMENT_DAY_woe	0.030390
4	MARITAL_STATUS_woe	0.027959
5	ACCOMPLISHMENT_woe	0.021630
2	OCCUPATION_TYPE_woe	0.018840
1	RESIDENCIAL_STATE_woe	0.016862
0	PROFESSIONAL_ZIP_3_woe	0.015906

4 Define Credit Scorecard

```
[112]: # tring a set of number
score_sc = sc.scorecard(bins_final,score_logreg,train_woe.columns[1:
    ↪],points0=800,odds0=0.2,pdo=200)
my_score = sc.scorecard_ply(df3,score_sc,print_step=False)
my_score.describe()
```

```
[112]:
```

	score
count	49543.000000
mean	344.276043
std	97.534830
min	71.000000

25%	275.000000
50%	342.000000
75%	414.000000
max	632.000000

The statistics of strategy related to the score card can be found in the report

5 Random Forest

```
[113]: # one hot encode every categorical variable
df2["PROFESSIONAL_ZIP_3"] = df2["PROFESSIONAL_ZIP_3"].astype("category")
df2["RESIDENCIAL_ZIP_3"] = df2["RESIDENCIAL_ZIP_3"].astype("category")
df2["PRODUCT"] = df2["PRODUCT"].astype("category")
df2["OCCUPATION_TYPE"] = df2["OCCUPATION_TYPE"].astype("category")
df2["PROFESSION_CODE"] = df2["PROFESSION_CODE"].astype("category")
df2["FLAG_PROFESSIONAL_PHONE"] = df2["FLAG_PROFESSIONAL_PHONE"].
    ↪astype("category")
df2["COMPANY"] = df2["COMPANY"].astype("category")
df2["QUANT_SPECIAL_BANKING_ACCOUNTS"] = df2["QUANT_SPECIAL_BANKING_ACCOUNTS"].
    ↪astype("category")
df2["QUANT_BANKING_ACCOUNTS"] = df2["QUANT_BANKING_ACCOUNTS"].astype("category")
df2["FLAG_OTHER_CARDS"] = df2["FLAG_OTHER_CARDS"].astype("category")
df2["FLAG_AMERICAN_EXPRESS"] = df2["FLAG_AMERICAN_EXPRESS"].astype("category")
df2["FLAG_DINERS"] = df2["FLAG_DINERS"].astype("category")
df2["FLAG_MASTERCARD"] = df2["FLAG_MASTERCARD"].astype("category")
df2["FLAG_VISA"] = df2["FLAG_VISA"].astype("category")
df2["FLAG_EMAIL"] = df2["FLAG_EMAIL"].astype("category")
df2["RESIDENCE_TYPE"] = df2["RESIDENCE_TYPE"].astype("category")
df2["MARITAL_STATUS"] = df2["MARITAL_STATUS"].astype("category")
df2["RESIDENCIAL_STATE"] = df2["RESIDENCIAL_STATE"].astype("category")
df2["PROFESSIONAL_ZIP_3"] = df2["PROFESSIONAL_ZIP_3"].astype("category")
df2["APPLICATION_SUBMISSION_TYPE"] = df2["APPLICATION_SUBMISSION_TYPE"].
    ↪astype("category")
df2["POSTAL_ADDRESS_TYPE"] = df2["POSTAL_ADDRESS_TYPE"].astype("category")
df2["QUANT_DEPENDANTS"] = df2["QUANT_DEPENDANTS"].astype("category")
df2["STATE_OF_BIRTH"] = df2["STATE_OF_BIRTH"].astype("category")
df2["CITY_OF_BIRTH"] = df2["CITY_OF_BIRTH"].astype("category")
df2["NACIONALITY"] = df2["NACIONALITY"].astype("category")
df2["RESIDENCIAL_STATE"] = df2["RESIDENCIAL_STATE"].astype("category")
df2["RESIDENCIAL_CITY"] = df2["RESIDENCIAL_CITY"].astype("category")
df2["RESIDENCIAL_BOROUGH"] = df2["RESIDENCIAL_BOROUGH"].astype("category")
df2["FLAG_RESIDENCIAL_PHONE"] = df2["FLAG_RESIDENCIAL_PHONE"].astype("category")
df7 = pd.get_dummies(df2)
```

```
[114]: # split again
train, test = sc.split_df(df7,y="TARGET_LABEL_BAD=1",ratio=0.7,seed=250918939).
    ↪values()
```

```
[117]: # define random forest classifier
score_rf = RandomForestClassifier(n_estimators = 1000,
                                criterion = "entropy",
                                max_depth=None,
                                min_samples_split=2,
                                min_samples_leaf=0.0001,
                                min_weight_fraction_leaf = 0,
                                max_features = "auto",
                                max_leaf_nodes=None,
                                min_impurity_decrease=0.0001,
                                bootstrap = True,
                                oob_score = True,
                                n_jobs=6,
                                random_state = 250918939,
                                verbose = 1,
                                warm_start = False,
                                class_weight="balanced"
                                )
```

```
[118]: # Train the RF
score_rf.fit(train.
    ↳drop(axis=1,columns="TARGET_LABEL_BAD=1"),train["TARGET_LABEL_BAD=1"])
```

```
[Parallel(n_jobs=6)]: Using backend ThreadingBackend with 6 concurrent workers.
[Parallel(n_jobs=6)]: Done 38 tasks      | elapsed:    1.6s
[Parallel(n_jobs=6)]: Done 188 tasks     | elapsed:    7.5s
[Parallel(n_jobs=6)]: Done 438 tasks     | elapsed:   17.3s
[Parallel(n_jobs=6)]: Done 788 tasks     | elapsed:   31.2s
[Parallel(n_jobs=6)]: Done 1000 out of 1000 | elapsed:   39.6s finished
W:\Tools\Anaconda3\envs\gpu\lib\site-packages\sklearn\base.py:445: UserWarning:
X does not have valid feature names, but RandomForestClassifier was fitted with
feature names
    warnings.warn(
```

```
[118]: RandomForestClassifier(class_weight='balanced', criterion='entropy',
                             min_impurity_decrease=0.0001, min_samples_leaf=0.0001,
                             min_weight_fraction_leaf=0, n_estimators=1000, n_jobs=6,
                             oob_score=True, random_state=250918939, verbose=1)
```

5.1 check performance

```
[119]: # Calculate confusion matrix
confusion_matrix_rf = confusion_matrix(y_true=test["TARGET_LABEL_BAD=1"],y_pred_
    ↳= pred_test)

# turn matrix to percentage
confusion_matrix_rf = confusion_matrix_rf.astype("float") / confusion_matrix_rf.
    ↳sum(axis=1)[:,np.newaxis]
```

```

# turn to dataframe
df_cm = pd.DataFrame(confusion_matrix_rf, index=["good", "bad"],
    ↪columns=["good", "bad"])

# parameters of the image
figsize = (10, 7)
fontsize = 14


# Create image
fig = plt.figure(figsize=figsize)
heatmap = sns.heatmap(df_cm, annot=True, fmt=".2f")

heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0,
    ↪ha="right", fontsize=fontsize)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45,
    ↪ha="right", fontsize=fontsize)

plt.ylabel("True label")
plt.xlabel("Predicted label")

plt.show()

```

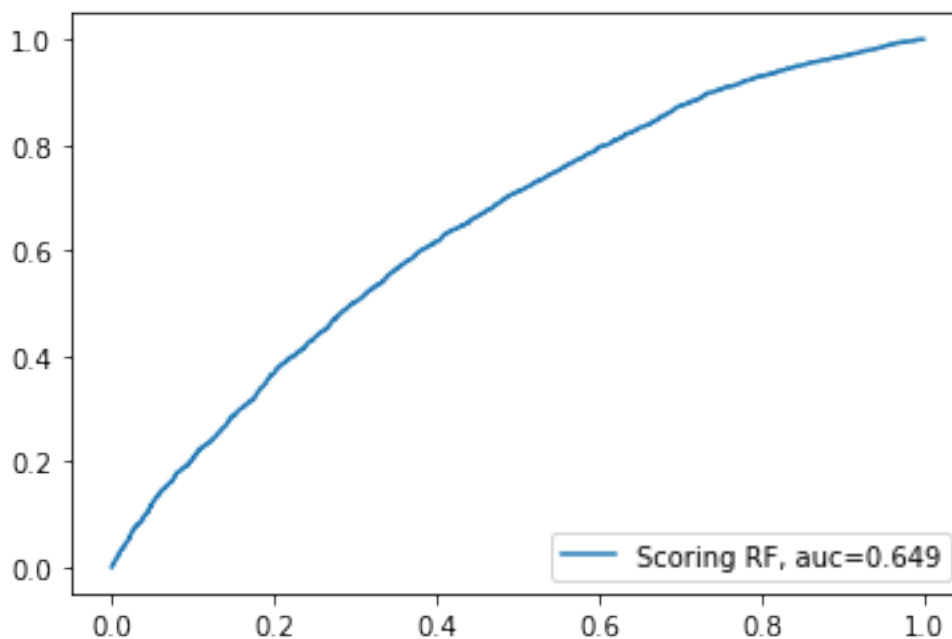


```
[120]: # calculate the ROC curve points
pred_test_prob = score_rf.predict_proba(test.
↳drop(axis=1,columns="TARGET_LABEL_BAD=1"))
fpr, tpr, thresholds = roc_curve(test["TARGET_LABEL_BAD=1"],pred_test_prob[:,1])

auc = np.round(roc_auc_score(y_true=test["TARGET_LABEL_BAD=1"],
↳y_score=pred_test_prob[:,1]),decimals=3)

plt.plot(fpr, tpr, label="Scoring RF, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```

```
[Parallel(n_jobs=6)]: Using backend ThreadingBackend with 6 concurrent workers.
[Parallel(n_jobs=6)]: Done 38 tasks      | elapsed:    0.0s
[Parallel(n_jobs=6)]: Done 188 tasks     | elapsed:    0.0s
[Parallel(n_jobs=6)]: Done 438 tasks     | elapsed:    0.2s
[Parallel(n_jobs=6)]: Done 788 tasks     | elapsed:    0.4s
[Parallel(n_jobs=6)]: Done 1000 out of 1000 | elapsed:    0.5s finished
```




5.2 variable importance

```
[121]: # variable importance
importances = score_rf.feature_importances_
indices = np.argsort(importances)[::-1]
columns = train.drop(axis=1, columns="TARGET_LABEL_BAD=1")
columns = columns.columns
columns_importance = columns[indices]
#print(columns_importance[:30])

for i in np.arange(13):
    print("{0:<30s}{1}".
        ↪format(columns_importance[i], round(importances[indices[i]], 4)))
```

AGE	0.0885
ACCOMPLISHMENT	0.0487
PAYMENT_DAY	0.0334
TOTAL_INCOME	0.0284
PERSONAL_MONTHLY_INCOME	0.0281
SUPPORT_POWER	0.0277
FLAG_RESIDENCIAL_PHONE_N	0.025
FLAG_RESIDENCIAL_PHONE_Y	0.0237
MONTHS_IN_RESIDENCE	0.0183
MARITAL_STATUS_1	0.0182
OCCUPATION_TYPE_1.0	0.0141
MARITAL_STATUS_2	0.0097
CITY_OF_BIRTH_OTHERS	0.0085

```
[122]: # plot variable importance
f, ax = plt.subplots(figsize=(3,8))
plt.title("Variable Importance - Random Forest")
sns.set_color_codes("pastel")
train2 = train.drop(axis=1, columns="TARGET_LABEL_BAD=1")
sns.barplot(y=[train2.columns[i] for i in indices[:15]], x=importances[indices[:
    ↪15]],
            label="Total", color="b")
ax.set(ylabel="Variable",
        xlabel="Variable Importance (Entropy)")
sns.despine(left=True, bottom=True)
```



6 XGBoosting

```
[144]: # define the classifier
score_XGB = XGBClassifier(max_depth=2, learning_rate=0.
    ↳1,n_estimators=50,verbosity=1,objective="binary:logistic",
```

```

        booster="gbtree",n_jobs=5,gamma=0.01,subsample=0.
↪632,colsample_bytree=1,colsample_bylevel=1,
        ↵
↪colsample_bynode=1,reg_alpha=1,reg_lambda=0,scale_pos_weight=36592/12951,
        base_score=0.
↪5,random_state=250918939,missing=1,tree_method="gpu_hist",
        gpu_id=1)

```

6.1 use cross validation to find optimal parameters

```

[145]: # potential parameters
param_grid=dict({"n_estimators": [50,75,100],
                "max_depth": [2,3,4],
                "learning_rate": [0.01,0.05,0.1,0.15]})

# reduce the sample to train for the parameters
val_train = train.sample(frac=0.2, random_state=250918939)

# define the grid search object
GridXGB = GridSearchCV(score_XGB, param_grid, cv=5, scoring="roc_auc", ↵
↪n_jobs=5, refit=False, verbose=2)
# train for the optimal parameter
GridXGB.fit(val_train.
↪drop(axis=1,columns="TARGET_LABEL_BAD=1"),val_train["TARGET_LABEL_BAD=1"])

```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

```

[145]: GridSearchCV(cv=5,
                estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                                        colsample_bylevel=1, colsample_bynode=1,
                                        colsample_bytree=1,
                                        enable_categorical=False, gamma=0.01,
                                        gpu_id=1, importance_type=None,
                                        interaction_constraints=None,
                                        learning_rate=0.1, max_delta_step=None,
                                        max_depth=2, min_child_weight=None,
                                        missing=1, monotone_constraints=None,
                                        n_estimators=50...
                                        num_parallel_tree=None, predictor=None,
                                        random_state=250918939, reg_alpha=1,
                                        reg_lambda=0, scale_pos_weight=None,
                                        scale_pos_weight=2.8254188865724656,
                                        subsample=0.632, tree_method='gpu_hist',
                                        validate_parameters=None, verbosity=1),
                n_jobs=5,
                param_grid={'learning_rate': [0.01, 0.05, 0.1, 0.15],
                            'max_depth': [2, 3, 4],

```



```
        'n_estimators': [50, 75, 100]],
    refit=False, scoring='roc_auc', verbose=2)
```

```
[146]: # show best params
print("The best AUC is %.3f" % GridXGB.best_score_)
GridXGB.best_params_
```

The best AUC is 0.621

```
[146]: {'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 75}
```

6.2 fit the model with best parameters

```
[149]: # create XG boosting with the best parameter
score_XGB = XGBClassifier(max_depth=2, learning_rate=0.
    ↳1,n_estimators=100,verbosity=1,objective="binary:logistic",
        booster="gbtree",n_jobs=2,gamma=0.000001,subsample=0.
    ↳632,colsample_bytree=1,colsample_bylevel=1,
        ↳
    ↳colsample_bynode=1,reg_alpha=1,reg_lambda=0,scale_pos_wieght=36592/12951,
        base_score=0.
    ↳5,random_state=250918939,missing=1,tree_method="gpu_hist",
        gpu_id=0)

# train over all training data
score_XGB.fit(train.
    ↳drop(axis=1,columns="TARGET_LABEL_BAD=1"),train["TARGET_LABEL_BAD=1"])
```

W:\Tools\Anaconda3\envs\gpu\lib\site-packages\xgboost\sklearn.py:1224:

UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[00:52:23] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:576:
Parameters: { "scale_pos_wieght" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but

then being mistakenly passed down to XGBoost core, or some parameter actually being used

but getting flagged wrongly here. Please open an issue if you find any such cases.

[00:52:24] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
[149]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
                    gamma=1e-06, gpu_id=0, importance_type=None,
                    interaction_constraints='', learning_rate=0.1, max_delta_step=0,
                    max_depth=2, min_child_weight=1, missing=1,
                    monotone_constraints='()', n_estimators=100, n_jobs=2,
                    num_parallel_tree=1, predictor='auto', random_state=250918939,
                    reg_alpha=1, reg_lambda=0, scale_pos_weight=1,
                    scale_pos_weight=2.8254188865724656, subsample=0.632,
                    tree_method='gpu_hist', validate_parameters=1, verbosity=1)
```

6.3 test performance

```
[150]: # apply the model to test set
pred_test2 = score_XGB.predict(test.drop(axis=1, columns="TARGET_LABEL_BAD=1"))

# Calculate confusion matrix
confusion_matrix_XGB =
    ↪ confusion_matrix(y_true=test["TARGET_LABEL_BAD=1"], y_pred = pred_test2)

# turn matrix to percentage
confusion_matrix_XGB = confusion_matrix_XGB.astype("float") /
    ↪ confusion_matrix_XGB.sum(axis=1)[:, np.newaxis]

# turn to dataframe
df_cm = pd.DataFrame(confusion_matrix_XGB, index=["good", "bad"],
    ↪ columns=["good", "bad"])


# parameters of the image
figsize = (10, 7)
fontsize = 14

# Create image
fig = plt.figure(figsize=figsize)
heatmap = sns.heatmap(df_cm, annot=True, fmt=".2f")

heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0,
    ↪ ha="right", fontsize=fontsize)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45,
    ↪ ha="right", fontsize=fontsize)
```

```
plt.ylabel("True label")
plt.xlabel("Predicted label")

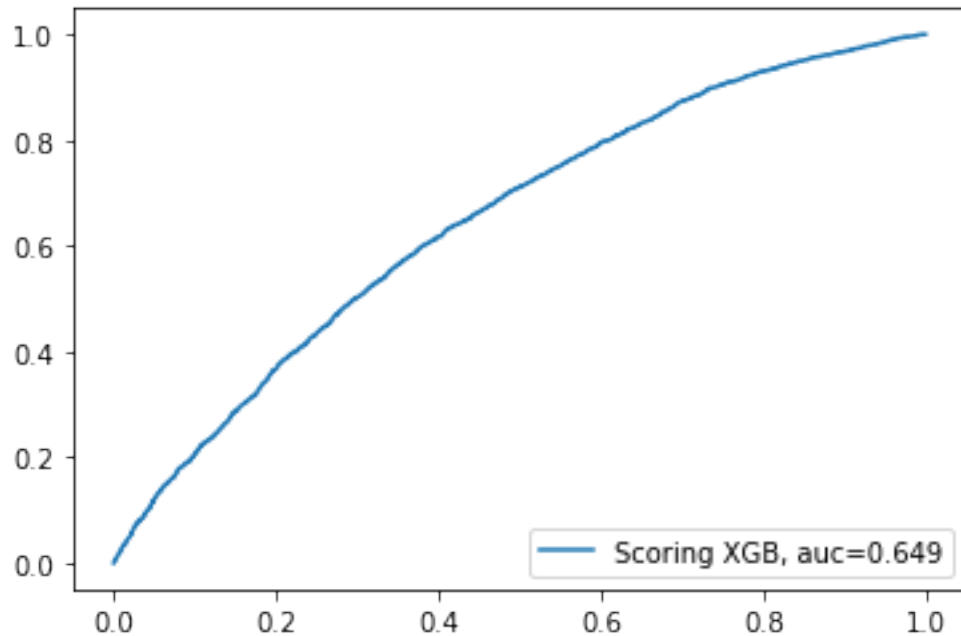
plt.show()
```



```
[151]: # calculate the ROC curve points
pred_test_prob2 = score_XGB.predict_proba(test.
    ↳ drop(axis=1, columns="TARGET_LABEL_BAD=1"))
fpr, tpr, thresholds = roc_curve(test["TARGET_LABEL_BAD=1"], pred_test_prob[:, 1])

auc = np.round(roc_auc_score(y_true=test["TARGET_LABEL_BAD=1"],
    ↳ y_score=pred_test_prob[:, 1]), decimals=3)

plt.plot(fpr, tpr, label="Scoring XGB, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



6.4 variable importance

```
[152]: # plot variable importance
importances2 = score_XGB.feature_importances_
indices2 = np.argsort(importances2)[::-1]

f, ax = plt.subplots(figsize=(3,8))
plt.title("Variable Importance - XGBoosting")
sns.set_color_codes("pastel")
train2 = train.drop(axis=1, columns="TARGET_LABEL_BAD=1")
sns.barplot(y=[train2.columns[i] for i in indices2[:15]],
            x=importances2[indices2[:15]],
            label="Total", color="b")
ax.set(ylabel="Variable",
       xlabel="Variable Importance (Entropy)")
sns.despine(left=True, bottom=True)
```

