

Assignment 4

March 30, 2022

```
[1]: # imports
import numpy as np
import matplotlib.pyplot as plt
import torch
import torchvision
import torch.nn as nn
import torch.nn.functional as F
import torchvision.transforms as transforms
import torchvision.models as models
from torch.utils.data import DataLoader
from torch import optim
from tqdm.auto import tqdm
import math
torch.manual_seed(55)
```

```
[1]: <torch._C.Generator at 0x286494ef510>
```

```
[ ]:
```

1 read in data

```
[2]: path = "W://Study Material/Jupyter Notebook/Datasets/CIFAR-10/"
      ↪ "cifar-10-batches-py"
def unpickle(file):
    import pickle
    with open(file, "rb") as fo:
        dict = pickle.load(fo, encoding="bytes")
    return dict
```

```
[3]: training_batch = np.empty((4), dtype=object)
      val_batch = np.empty((1), dtype=object)
      test_batch = np.empty((1), dtype=object)
      training_batch.shape
```

```
[3]: (4,)
```

```
[4]: training_batch[0] = unpickle(path+"/data_batch_1")
training_batch[1] = unpickle(path+"/data_batch_2")
training_batch[2] = unpickle(path+"/data_batch_3")
training_batch[3] = unpickle(path+"/data_batch_4")
val_batch[0] = unpickle(path+"/data_batch_5")
test_batch[0] = unpickle(path+"/test_batch")
meta_data = unpickle(path+"/batches.meta")

[5]: training_batch[0].keys()

[5]: dict_keys([b'batch_label', b'labels', b'data', b'filenames'])

[6]: np.array(training_batch[0][b'labels']).shape

[6]: (10000,)

[7]: training_batch[0][b'data'].shape

[7]: (10000, 3072)

[8]: meta_data[b'label_names']

[8]: [b'airplane',
      b'automobile',
      b'bird',
      b'cat',
      b'deer',
      b'dog',
      b'frog',
      b'horse',
      b'ship',
      b'truck']

[ ]:
```

2 define custom data loader

```
[9]: class CustomDataset(torch.utils.data.Dataset):
      def __init__(self, batch, b_type, transform):
          self.transform = transform
          if b_type == "train":
              self.data = np.concatenate((batch[0][b'data'], batch[1][b'data'],
              ↵
              ↪ batch[2][b'data'], batch[3][b'data']), axis=0)
              #print(b_type, " data shape: ", self.data.shape)
              self.label = np.
              ↪ concatenate((batch[0][b'labels'], batch[1][b'labels'],
```

```

        ↪ batch[2][b'labels'], batch[3][b'labels']))
        #print(b_type, " label shape: ", len(self.label))
    else:
        self.data = batch[0][b'data']
        #print(b_type, " data shape: ", self.data.shape)
        self.label = batch[0][b'labels']
        #print(b_type, " label shape: ", len(self.label))

    def __len__(self):
        return len(self.label)

    def __getitem__(self, index):
        label = self.label[index]
        image_data = self.data[index]
        # convert (3072,) array to (3,32,32)
        image_r = image_data[:1024].reshape(32,32)
        image_g = image_data[1024:2048].reshape(32,32)
        image_b = image_data[2048:].reshape(32,32)
        image = np.array([image_r, image_g, image_b])
        image = np.transpose(image, (1,2,0)) # change the batch dimension to the ↪
        ↪ last dimension
        #image = torch.tensor(image, dtype=torch.float)
        if self.transform is not None:
            image = self.transform(image)
        #print(image.shape)
        return image, label

```

```

[10]: x = np.ones((1,2,3))
      np.transpose(x, (2,0,1)).shape

```

```

[10]: (3, 1, 2)

```

2.0.1 transformation

```

[11]: # define data augmentation
data_transformers = {"train": transforms.Compose([transforms.ToPILImage(),
                                                  transforms.RandomResizedCrop(256),
                                                  transforms.RandomHorizontalFlip(),
                                                  transforms.RandomRotation(30),
                                                  transforms.ToTensor(),
                                                  transforms.Normalize((0.5,0.5,0.5),(0.5,0.
        ↪ 5,0.5)))],
                    "test": transforms.Compose([transforms.ToPILImage(),
                                                  transforms.Resize(256),
                                                  transforms.ToTensor(),

```

```
transforms.Normalize((0.5,0.5,0.5),(0.5,0.5,0.5))]]}
```

2.0.2 datasets

```
[12]: # defining datasets
train_data = CustomDataset(training_batch,"train",transform=data_transformers["train"])
val_data = CustomDataset(val_batch,"validation",transform=data_transformers["test"])
test_data = CustomDataset(test_batch,"test",transform=data_transformers["test"])
```

```
[13]: # make sure the shape of the image extracted is okay
print(train_data[0][0].shape)
#print(train_data[0][1])
```

```
torch.Size([3, 256, 256])
```

2.0.3 dataloader

```
[14]: train_dataloader = DataLoader(train_data,batch_size=128,shuffle=True)
val_dataloader = DataLoader(val_data,batch_size=128,shuffle=True)
test_dataloader = DataLoader(test_data,batch_size=128,shuffle=False)
```

```
[ ]:
```

```
[ ]:
```

3 training and validation routine

```
[20]: def
    train(model,trainloader,valloader,epochs,print_frequency,loss_fn,optimizer,device,run,part)

    training_steps = epochs * (len(trainloader))
    progress_bar = tqdm(range(training_steps))
    ep = epochs
    print_every = math.floor(training_steps/print_frequency)
    steps = 0
    model.to(device)
    model.train()
    total = 0
    correct = 0
    best_val = 10000

    current_loss = 0

    for e in np.arange(ep):
```

```

for batch,(images,labels) in enumerate(trainloader):
    steps += 1
    images, labels = images.to(device), labels.to(device)
    # convert images to float because weights are floats
    images = images.float()
    labels = labels.type(torch.long) # need to be int

    # calculate loss and backpropagate
    optimizer.zero_grad()
    outputs = model(images)
    _,predictions = torch.max(outputs.data,1)
    total += labels.size(0)
    correct += (predictions == labels).sum().item()
    loss = loss_fn(outputs,labels)
    loss.backward()
    optimizer.step()

    current_loss += loss.item()
    progress_bar.update(1)

    if steps % print_every == 0:
        print('EPOCHS : {}/{}'.format(e+1,epochs),
              'Loss : {:.6f}'.format(current_loss/print_every))
        print('The training accuracy is {:.2f}%'.format(correct/
→total*100))

        current_loss = 0
        val_loss = validate(model,valloader,loss_fn, device)
        if val_loss < best_val:
            torch.save(model.state_dict(),"Weights/part-{}-run-{}-Best.
→pth".format(str(part),

→str(run)))
            torch.save(model.state_dict(),"Weights/part-{}-run-{}-Last.pth".
→format(str(part),

→str(run)))

```

```

[21]: def validate(model, valloader, loss_fn, device):
    total = 0
    correct = 0
    val_loss = 0
    model.eval()
    steps = 0
    with torch.no_grad():
        for batch, (images,labels) in enumerate(valloader):
            images, labels = images.to(device), labels.to(device)
            # convert images to float because weights are floats

```

```

        images = images.float()
        labels = labels.type(torch.long) # need to be int
        steps += 1
        outputs = model(images)
        _, predictions = torch.max(outputs.data,1)
        total += labels.size(0)
        correct += (labels == predictions).sum().item()
        loss = loss_fn(outputs,labels)
        val_loss += loss.item()
    val_loss /= steps
    accuracy = correct / total * 100
    print("The validation loss is %.4f" % (val_loss))
    print('The valudation accuracy is {:.2f}%\n'.format(accuracy))
    return val_loss

```

[]:

[]:

[]:

[]:

[]:

4 Part b)

```

[17]: AlexNet = models.alexnet(pretrained=False)
AlexNet.load_state_dict(torch.load("W://Study Material/Jupyter Notebook/
↳Pretrained_Weights/alexnet-owt-7be5be79.pth"))
AlexNet.cuda()

```

```

[17]: AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)

```

```

        (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (11): ReLU(inplace=True)
        (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
    (classifier): Sequential(
      (0): Dropout(p=0.5, inplace=False)
      (1): Linear(in_features=9216, out_features=4096, bias=True)
      (2): ReLU(inplace=True)
      (3): Dropout(p=0.5, inplace=False)
      (4): Linear(in_features=4096, out_features=4096, bias=True)
      (5): ReLU(inplace=True)
      (6): Linear(in_features=4096, out_features=1000, bias=True)
    )
  )

```

```
[ ]:
```

```
[18]: loss_fn = nn.CrossEntropyLoss()
optimizer = optim.Adam(AlexNet.parameters(),lr=1e-3)
```

```
[24]: train(AlexNet,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",1)
```

```
0%|          | 0/3130 [00:00<?, ?it/s]
```

```

EPOCHS : 1/10 Loss : 2.247014
The training accuracy is 21.80%
The validation loss is 1.6998
The valudation accuracy is 36.71%

```

```

EPOCHS : 1/10 Loss : 1.785275
The training accuracy is 27.76%
The validation loss is 1.5200
The valudation accuracy is 43.46%

```

```

EPOCHS : 2/10 Loss : 1.685747
The training accuracy is 31.12%
The validation loss is 1.4152
The valudation accuracy is 47.11%

```

```

EPOCHS : 2/10 Loss : 1.600968
The training accuracy is 33.61%
The validation loss is 1.3193
The valudation accuracy is 51.00%

```

```

EPOCHS : 3/10 Loss : 1.554594
The training accuracy is 35.58%

```

The validation loss is 1.3243
The valudation accuracy is 53.44%

EPOCHS : 3/10 Loss : 1.490266
The training accuracy is 37.28%
The validation loss is 1.1996
The valudation accuracy is 56.95%

EPOCHS : 4/10 Loss : 1.486563
The training accuracy is 38.56%
The validation loss is 1.1369
The valudation accuracy is 58.76%

EPOCHS : 4/10 Loss : 1.429415
The training accuracy is 39.73%
The validation loss is 1.1196
The valudation accuracy is 60.99%

EPOCHS : 5/10 Loss : 1.401546
The training accuracy is 40.79%
The validation loss is 1.0909
The valudation accuracy is 61.05%

EPOCHS : 5/10 Loss : 1.368858
The training accuracy is 41.79%
The validation loss is 1.0646
The valudation accuracy is 62.20%

EPOCHS : 6/10 Loss : 1.354280
The training accuracy is 42.69%
The validation loss is 1.0736
The valudation accuracy is 62.41%

EPOCHS : 6/10 Loss : 1.334924
The training accuracy is 43.50%
The validation loss is 0.9667
The valudation accuracy is 65.68%

EPOCHS : 7/10 Loss : 1.300909
The training accuracy is 44.25%
The validation loss is 0.9733
The valudation accuracy is 66.59%

EPOCHS : 7/10 Loss : 1.288129
The training accuracy is 44.92%
The validation loss is 1.0068
The valudation accuracy is 64.91%

EPOCHS : 8/10 Loss : 1.270171
The training accuracy is 45.57%
The validation loss is 0.9393
The valudation accuracy is 67.63%

EPOCHS : 8/10 Loss : 1.249892
The training accuracy is 46.21%
The validation loss is 0.9625
The valudation accuracy is 67.05%

EPOCHS : 9/10 Loss : 1.238430
The training accuracy is 46.80%
The validation loss is 0.8542
The valudation accuracy is 70.39%

EPOCHS : 9/10 Loss : 1.230177
The training accuracy is 47.33%
The validation loss is 0.8584
The valudation accuracy is 70.26%

EPOCHS : 10/10 Loss : 1.208205
The training accuracy is 47.83%
The validation loss is 0.8575
The valudation accuracy is 70.58%

EPOCHS : 10/10 Loss : 1.206469
The training accuracy is 48.30%
The validation loss is 0.8517
The valudation accuracy is 70.97%

```
[25]: optimizer = optim.Adam(AlexNet.parameters(),lr=1e-5)
      train(AlexNet,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",2)
```

```
0%|          | 0/3130 [00:00<?, ?it/s]
```

EPOCHS : 1/10 Loss : 1.314707
The training accuracy is 54.08%
The validation loss is 0.7874
The valudation accuracy is 72.85%

EPOCHS : 1/10 Loss : 1.110474
The training accuracy is 57.41%
The validation loss is 0.7677
The valudation accuracy is 73.35%

EPOCHS : 2/10 Loss : 1.099722
The training accuracy is 58.54%
The validation loss is 0.7579

The valudation accuracy is 73.52%

EPOCHS : 2/10 Loss : 1.087660

The training accuracy is 59.29%

The validation loss is 0.7448

The valudation accuracy is 74.05%

EPOCHS : 3/10 Loss : 1.072100

The training accuracy is 59.91%

The validation loss is 0.7409

The valudation accuracy is 74.36%

EPOCHS : 3/10 Loss : 1.068306

The training accuracy is 60.23%

The validation loss is 0.7364

The valudation accuracy is 74.53%

EPOCHS : 4/10 Loss : 1.055259

The training accuracy is 60.57%

The validation loss is 0.7403

The valudation accuracy is 74.46%

EPOCHS : 4/10 Loss : 1.051216

The training accuracy is 60.83%

The validation loss is 0.7277

The valudation accuracy is 74.69%

EPOCHS : 5/10 Loss : 1.054497

The training accuracy is 61.01%

The validation loss is 0.7298

The valudation accuracy is 74.64%

EPOCHS : 5/10 Loss : 1.037197

The training accuracy is 61.21%

The validation loss is 0.7229

The valudation accuracy is 74.84%

EPOCHS : 6/10 Loss : 1.044361

The training accuracy is 61.40%

The validation loss is 0.7249

The valudation accuracy is 75.03%

EPOCHS : 6/10 Loss : 1.044632

The training accuracy is 61.51%

The validation loss is 0.7188

The valudation accuracy is 74.97%

EPOCHS : 7/10 Loss : 1.044561

The training accuracy is 61.61%
The validation loss is 0.7205
The valudation accuracy is 75.01%

EPOCHS : 7/10 Loss : 1.030355
The training accuracy is 61.74%
The validation loss is 0.7082
The valudation accuracy is 75.07%

EPOCHS : 8/10 Loss : 1.045026
The training accuracy is 61.83%
The validation loss is 0.7083
The valudation accuracy is 75.31%

EPOCHS : 8/10 Loss : 1.014952
The training accuracy is 61.98%
The validation loss is 0.7143
The valudation accuracy is 75.18%

EPOCHS : 9/10 Loss : 1.038396
The training accuracy is 62.04%
The validation loss is 0.7163
The valudation accuracy is 75.08%

EPOCHS : 9/10 Loss : 1.020153
The training accuracy is 62.15%
The validation loss is 0.7034
The valudation accuracy is 75.28%

EPOCHS : 10/10 Loss : 1.028688
The training accuracy is 62.22%
The validation loss is 0.7066
The valudation accuracy is 75.30%

EPOCHS : 10/10 Loss : 1.011283
The training accuracy is 62.33%
The validation loss is 0.6973
The valudation accuracy is 75.38%

```
[19]: optimizer = optim.Adam(AlexNet.parameters(),lr=1e-6)
      AlexNet.load_state_dict(torch.load("Weights/run-2-Best.pth"))
      train(AlexNet,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",3)
```

```
0%|          | 0/3130 [00:00<?, ?it/s]
```

EPOCHS : 1/10 Loss : 1.236205
The training accuracy is 57.47%
The validation loss is 0.7003

The valudation accuracy is 75.54%

EPOCHS : 1/10 Loss : 1.020405

The training accuracy is 60.70%

The validation loss is 0.7034

The valudation accuracy is 75.55%

EPOCHS : 2/10 Loss : 1.018962

The training accuracy is 61.69%

The validation loss is 0.6994

The valudation accuracy is 75.49%

EPOCHS : 2/10 Loss : 1.014674

The training accuracy is 62.28%

The validation loss is 0.6986

The valudation accuracy is 75.52%

EPOCHS : 3/10 Loss : 1.023155

The training accuracy is 62.52%

The validation loss is 0.7007

The valudation accuracy is 75.56%

EPOCHS : 3/10 Loss : 1.013412

The training accuracy is 62.75%

The validation loss is 0.7012

The valudation accuracy is 75.43%

EPOCHS : 4/10 Loss : 1.010199

The training accuracy is 62.95%

The validation loss is 0.7021

The valudation accuracy is 75.45%

EPOCHS : 4/10 Loss : 1.028888

The training accuracy is 63.05%

The validation loss is 0.7022

The valudation accuracy is 75.49%

EPOCHS : 5/10 Loss : 1.016358

The training accuracy is 63.15%

The validation loss is 0.6972

The valudation accuracy is 75.48%

EPOCHS : 5/10 Loss : 1.023578

The training accuracy is 63.22%

The validation loss is 0.7030

The valudation accuracy is 75.59%

EPOCHS : 6/10 Loss : 1.019098

The training accuracy is 63.26%
The validation loss is 0.7011
The valudation accuracy is 75.58%

EPOCHS : 6/10 Loss : 1.015666
The training accuracy is 63.36%
The validation loss is 0.7017
The valudation accuracy is 75.58%

EPOCHS : 7/10 Loss : 1.015819
The training accuracy is 63.42%
The validation loss is 0.7019
The valudation accuracy is 75.66%

EPOCHS : 7/10 Loss : 0.999843
The training accuracy is 63.52%
The validation loss is 0.6974
The valudation accuracy is 75.63%

EPOCHS : 8/10 Loss : 1.012752
The training accuracy is 63.54%
The validation loss is 0.7003
The valudation accuracy is 75.61%

EPOCHS : 8/10 Loss : 1.019031
The training accuracy is 63.54%
The validation loss is 0.7079
The valudation accuracy is 75.57%

EPOCHS : 9/10 Loss : 1.016139
The training accuracy is 63.54%
The validation loss is 0.7064
The valudation accuracy is 75.57%

EPOCHS : 9/10 Loss : 1.005403
The training accuracy is 63.60%
The validation loss is 0.6966
The valudation accuracy is 75.55%

EPOCHS : 10/10 Loss : 1.006781
The training accuracy is 63.62%
The validation loss is 0.6981
The valudation accuracy is 75.61%

EPOCHS : 10/10 Loss : 1.022852
The training accuracy is 63.63%
The validation loss is 0.6972
The valudation accuracy is 75.57%

[]:

[]:

[]:

[]:

[]:

5 Part c)

```
[26]: AlexNet2 = models.alexnet(pretrained=False)
AlexNet2.load_state_dict(torch.load("W://Study Material/Jupyter Notebook/
↳Pretrained_Weights/alexnet-owt-7be5be79.pth"))
AlexNet2.cuda()
```

```
[26]: AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
  ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
  ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
  ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
```

```
)
)
```

```
[27]: AlexNet2.classifier = nn.Sequential(
    nn.Dropout(0.5,False),
    nn.Linear(9216,4096,True),
    nn.ReLU(True),
    nn.Linear(4096,10,True)
)
print(AlexNet2)
```

```
AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Linear(in_features=4096, out_features=10, bias=True)
  )
)
```

```
[24]: loss_fn = nn.CrossEntropyLoss()
optimizer = optim.Adam(AlexNet2.parameters(),lr=1e-3)
```

```
[25]: train(AlexNet2,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",1,"c")
```

```
0%|          | 0/3130 [00:00<?, ?it/s]
```

```
EPOCHS : 1/10 Loss : 2.342232
```

```
The training accuracy is 9.85%
```

```
The validation loss is 2.3030
```

The valudation accuracy is 9.52%

EPOCHS : 1/10 Loss : 2.302850

The training accuracy is 9.81%

The validation loss is 2.3027

The valudation accuracy is 10.14%

EPOCHS : 2/10 Loss : 2.302646

The training accuracy is 9.81%

The validation loss is 2.3028

The valudation accuracy is 10.14%

EPOCHS : 2/10 Loss : 2.302732

The training accuracy is 9.78%

The validation loss is 2.3027

The valudation accuracy is 10.03%

```
-----  
KeyboardInterrupt                                Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_23084\577002224.py in <module>  
----> 1  
      ↪ train(AlexNet2, train_dataloader, val_dataloader, 10, 20, loss_fn, optimizer, "cuda" 1, "c")  
  
~\AppData\Local\Temp\ipykernel_23084\1989990385.py in train(model, trainloader, valloader, epochs, print_frequency, loss_fn, optimizer, device, run, part)  
      ↪  
      16         for batch, (images, labels) in enumerate(trainloader):  
      17             steps += 1  
----> 18             images, labels = images.to(device), labels.to(device)  
      19             # convert images to float because weights are floats  
      20             images = images.float()  
  
KeyboardInterrupt:
```

```
[28]: loss_fn = nn.CrossEntropyLoss()  
      optimizer = optim.Adam(AlexNet2.parameters(), lr=1e-5)
```

```
[29]: train(AlexNet2, train_dataloader, val_dataloader, 10, 20, loss_fn, optimizer, "cuda", 1, "c")
```

```
0%|          | 0/3130 [00:00<?, ?it/s]
```

EPOCHS : 1/10 Loss : 1.749224

The training accuracy is 37.26%

The validation loss is 1.0611

The valudation accuracy is 62.49%

EPOCHS : 1/10 Loss : 1.344720

The training accuracy is 44.61%

The validation loss is 0.8322
The valudation accuracy is 71.62%

EPOCHS : 2/10 Loss : 1.248277
The training accuracy is 48.05%
The validation loss is 0.7669
The valudation accuracy is 73.68%

EPOCHS : 2/10 Loss : 1.182601
The training accuracy is 50.66%
The validation loss is 0.7867
The valudation accuracy is 72.40%

EPOCHS : 3/10 Loss : 1.149564
The training accuracy is 52.32%
The validation loss is 0.6671
The valudation accuracy is 76.95%

EPOCHS : 3/10 Loss : 1.132360
The training accuracy is 53.53%
The validation loss is 0.6508
The valudation accuracy is 77.58%

EPOCHS : 4/10 Loss : 1.089950
The training accuracy is 54.61%
The validation loss is 0.6348
The valudation accuracy is 78.25%

EPOCHS : 4/10 Loss : 1.067065
The training accuracy is 55.53%
The validation loss is 0.5904
The valudation accuracy is 79.56%

EPOCHS : 5/10 Loss : 1.058476
The training accuracy is 56.33%
The validation loss is 0.5837
The valudation accuracy is 79.80%

EPOCHS : 5/10 Loss : 1.039308
The training accuracy is 57.00%
The validation loss is 0.5653
The valudation accuracy is 80.63%

EPOCHS : 6/10 Loss : 1.022074
The training accuracy is 57.62%
The validation loss is 0.5610
The valudation accuracy is 80.57%

EPOCHS : 6/10 Loss : 1.004124
The training accuracy is 58.18%
The validation loss is 0.5485
The valudation accuracy is 80.86%

EPOCHS : 7/10 Loss : 0.996803
The training accuracy is 58.69%
The validation loss is 0.5404
The valudation accuracy is 81.06%

EPOCHS : 7/10 Loss : 0.982126
The training accuracy is 59.15%
The validation loss is 0.5449
The valudation accuracy is 81.12%

EPOCHS : 8/10 Loss : 0.975503
The training accuracy is 59.57%
The validation loss is 0.5282
The valudation accuracy is 81.34%

EPOCHS : 8/10 Loss : 0.971869
The training accuracy is 59.97%
The validation loss is 0.5269
The valudation accuracy is 81.63%

EPOCHS : 9/10 Loss : 0.954651
The training accuracy is 60.34%
The validation loss is 0.4988
The valudation accuracy is 83.03%

EPOCHS : 9/10 Loss : 0.958505
The training accuracy is 60.65%
The validation loss is 0.4952
The valudation accuracy is 82.98%

EPOCHS : 10/10 Loss : 0.949912
The training accuracy is 60.97%
The validation loss is 0.4980
The valudation accuracy is 82.53%

EPOCHS : 10/10 Loss : 0.922114
The training accuracy is 61.29%
The validation loss is 0.4982
The valudation accuracy is 82.74%

```
[30]: optimizer = optim.Adam(AlexNet2.parameters(),lr=1e-7)
      train(AlexNet2,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",2,"c")
```

```
0%|          | 0/3130 [00:00<?, ?it/s]
```

```
EPOCHS : 1/10 Loss : 0.983893
The training accuracy is 65.21%
The validation loss is 0.4936
The valudation accuracy is 83.12%
```

```
EPOCHS : 1/10 Loss : 0.904120
The training accuracy is 66.81%
The validation loss is 0.4828
The valudation accuracy is 83.32%
```

```
EPOCHS : 2/10 Loss : 0.906161
The training accuracy is 67.25%
The validation loss is 0.4769
The valudation accuracy is 83.34%
```

```
EPOCHS : 2/10 Loss : 0.900338
The training accuracy is 67.54%
The validation loss is 0.4763
The valudation accuracy is 83.48%
```

```
EPOCHS : 3/10 Loss : 0.921114
The training accuracy is 67.56%
The validation loss is 0.4737
The valudation accuracy is 83.46%
```

```
EPOCHS : 3/10 Loss : 0.895418
The training accuracy is 67.72%
The validation loss is 0.4758
The valudation accuracy is 83.38%
```

```
EPOCHS : 4/10 Loss : 0.905156
The training accuracy is 67.83%
The validation loss is 0.4768
The valudation accuracy is 83.53%
```

```
EPOCHS : 4/10 Loss : 0.887147
The training accuracy is 67.99%
The validation loss is 0.4795
The valudation accuracy is 83.49%
```

```
EPOCHS : 5/10 Loss : 0.894696
The training accuracy is 68.05%
The validation loss is 0.4749
```

The valudation accuracy is 83.54%

EPOCHS : 5/10 Loss : 0.898819

The training accuracy is 68.12%

The validation loss is 0.4786

The valudation accuracy is 83.61%

EPOCHS : 6/10 Loss : 0.897295

The training accuracy is 68.15%

The validation loss is 0.4721

The valudation accuracy is 83.72%

EPOCHS : 6/10 Loss : 0.907410

The training accuracy is 68.16%

The validation loss is 0.4769

The valudation accuracy is 83.68%

EPOCHS : 7/10 Loss : 0.901949

The training accuracy is 68.16%

The validation loss is 0.4778

The valudation accuracy is 83.59%

EPOCHS : 7/10 Loss : 0.892131

The training accuracy is 68.21%

The validation loss is 0.4749

The valudation accuracy is 83.58%

EPOCHS : 8/10 Loss : 0.893424

The training accuracy is 68.25%

The validation loss is 0.4760

The valudation accuracy is 83.66%

EPOCHS : 8/10 Loss : 0.907078

The training accuracy is 68.25%

The validation loss is 0.4715

The valudation accuracy is 83.68%

EPOCHS : 9/10 Loss : 0.907165

The training accuracy is 68.24%

The validation loss is 0.4700

The valudation accuracy is 83.64%

EPOCHS : 9/10 Loss : 0.888679

The training accuracy is 68.26%

The validation loss is 0.4740

The valudation accuracy is 83.61%

EPOCHS : 10/10 Loss : 0.898744

The training accuracy is 68.29%
The validation loss is 0.4734
The valudation accuracy is 83.69%

EPOCHS : 10/10 Loss : 0.893277
The training accuracy is 68.30%
The validation loss is 0.4694
The valudation accuracy is 83.58%

```
[31]: optimizer = optim.Adam(AlexNet2.parameters(),lr=1e-8)
      #AlexNet.load_state_dict(torch.load("Weights/part-c-run-2-Best.pth"))
      train(AlexNet2,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",3,"c")
```

0%| | 0/3130 [00:00<?, ?it/s]

EPOCHS : 1/10 Loss : 0.960729
The training accuracy is 65.96%
The validation loss is 0.4728
The valudation accuracy is 83.51%

EPOCHS : 1/10 Loss : 0.884539
The training accuracy is 67.56%
The validation loss is 0.4713
The valudation accuracy is 83.52%

EPOCHS : 2/10 Loss : 0.894978
The training accuracy is 67.94%
The validation loss is 0.4709
The valudation accuracy is 83.56%

EPOCHS : 2/10 Loss : 0.900989
The training accuracy is 68.03%
The validation loss is 0.4811
The valudation accuracy is 83.58%

EPOCHS : 3/10 Loss : 0.895619
The training accuracy is 68.20%
The validation loss is 0.4710
The valudation accuracy is 83.62%

EPOCHS : 3/10 Loss : 0.904552
The training accuracy is 68.23%
The validation loss is 0.4682
The valudation accuracy is 83.59%

EPOCHS : 4/10 Loss : 0.898023
The training accuracy is 68.27%
The validation loss is 0.4718

The valudation accuracy is 83.64%

EPOCHS : 4/10 Loss : 0.896598

The training accuracy is 68.29%

The validation loss is 0.4702

The valudation accuracy is 83.62%

EPOCHS : 5/10 Loss : 0.886057

The training accuracy is 68.36%

The validation loss is 0.4696

The valudation accuracy is 83.62%

EPOCHS : 5/10 Loss : 0.891967

The training accuracy is 68.42%

The validation loss is 0.4730

The valudation accuracy is 83.62%

EPOCHS : 6/10 Loss : 0.902300

The training accuracy is 68.40%

The validation loss is 0.4713

The valudation accuracy is 83.61%

EPOCHS : 6/10 Loss : 0.888475

The training accuracy is 68.45%

The validation loss is 0.4702

The valudation accuracy is 83.65%

EPOCHS : 7/10 Loss : 0.886854

The training accuracy is 68.51%

The validation loss is 0.4714

The valudation accuracy is 83.63%

EPOCHS : 7/10 Loss : 0.900989

The training accuracy is 68.50%

The validation loss is 0.4767

The valudation accuracy is 83.65%

EPOCHS : 8/10 Loss : 0.898051

The training accuracy is 68.50%

The validation loss is 0.4729

The valudation accuracy is 83.66%

EPOCHS : 8/10 Loss : 0.888562

The training accuracy is 68.51%

The validation loss is 0.4689

The valudation accuracy is 83.66%

EPOCHS : 9/10 Loss : 0.882809

The training accuracy is 68.55%
The validation loss is 0.4671
The valudation accuracy is 83.67%

EPOCHS : 9/10 Loss : 0.898687
The training accuracy is 68.53%
The validation loss is 0.4736
The valudation accuracy is 83.69%

EPOCHS : 10/10 Loss : 0.892271
The training accuracy is 68.54%
The validation loss is 0.4707
The valudation accuracy is 83.68%

EPOCHS : 10/10 Loss : 0.904148
The training accuracy is 68.52%
The validation loss is 0.4711
The valudation accuracy is 83.65%

[]:

[]:

[]:

[]:

6 Part d)

```
[32]: AlexNet3 = models.alexnet(pretrained=False)
AlexNet3.load_state_dict(torch.load("W://Study Material/Jupyter Notebook/
↳Pretrained_Weights/alexnet-owt-7be5be79.pth"))
AlexNet3.cuda()
```

```
[32]: AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
```

```

(8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(9): ReLU(inplace=True)
(10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace=True)
(12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
(classifier): Sequential(
  (0): Dropout(p=0.5, inplace=False)
  (1): Linear(in_features=9216, out_features=4096, bias=True)
  (2): ReLU(inplace=True)
  (3): Dropout(p=0.5, inplace=False)
  (4): Linear(in_features=4096, out_features=4096, bias=True)
  (5): ReLU(inplace=True)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)

```

```

[34]: AlexNet3.classifier[6] = nn.Linear(4096,10,True)
print(AlexNet3)

```

```

AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,
ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
  )
)

```



```

        (5): ReLU(inplace=True)
        (6): Linear(in_features=4096, out_features=10, bias=True)
    )
)

```

```

[35]: loss_fn = nn.CrossEntropyLoss()
      optimizer = optim.Adam(AlexNet3.parameters(),lr=1e-5)

```

```

[36]: train(AlexNet3,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",1,"d")

```

```

0%|          | 0/3130 [00:00<?, ?it/s]

```

```

EPOCHS : 1/10 Loss : 1.850561
The training accuracy is 32.47%
The validation loss is 1.2623
The valudation accuracy is 55.76%

```

```

EPOCHS : 1/10 Loss : 1.405315
The training accuracy is 40.75%
The validation loss is 0.9747
The valudation accuracy is 65.92%

```

```

EPOCHS : 2/10 Loss : 1.257076
The training accuracy is 45.55%
The validation loss is 0.8586
The valudation accuracy is 69.45%

```

```

EPOCHS : 2/10 Loss : 1.177570
The training accuracy is 48.66%
The validation loss is 0.7658
The valudation accuracy is 73.17%

```

```

EPOCHS : 3/10 Loss : 1.127938
The training accuracy is 50.88%
The validation loss is 0.6838
The valudation accuracy is 76.32%

```

```

EPOCHS : 3/10 Loss : 1.089887
The training accuracy is 52.55%
The validation loss is 0.6361
The valudation accuracy is 78.05%

```

```

EPOCHS : 4/10 Loss : 1.061739
The training accuracy is 53.94%
The validation loss is 0.6482
The valudation accuracy is 77.48%

```

```

EPOCHS : 4/10 Loss : 1.048316
The training accuracy is 55.06%

```

The validation loss is 0.5840
The valudation accuracy is 79.73%

EPOCHS : 5/10 Loss : 1.028207
The training accuracy is 56.01%
The validation loss is 0.5783
The valudation accuracy is 79.75%

EPOCHS : 5/10 Loss : 0.997779
The training accuracy is 56.87%
The validation loss is 0.5529
The valudation accuracy is 80.76%

EPOCHS : 6/10 Loss : 0.981257
The training accuracy is 57.63%
The validation loss is 0.5608
The valudation accuracy is 80.42%

EPOCHS : 6/10 Loss : 0.972149
The training accuracy is 58.28%
The validation loss is 0.5433
The valudation accuracy is 81.03%

EPOCHS : 7/10 Loss : 0.949519
The training accuracy is 58.90%
The validation loss is 0.5389
The valudation accuracy is 81.30%

EPOCHS : 7/10 Loss : 0.933250
The training accuracy is 59.49%
The validation loss is 0.5344
The valudation accuracy is 81.23%

EPOCHS : 8/10 Loss : 0.917924
The training accuracy is 60.05%
The validation loss is 0.5027
The valudation accuracy is 82.17%

EPOCHS : 8/10 Loss : 0.908061
The training accuracy is 60.55%
The validation loss is 0.4832
The valudation accuracy is 83.11%

EPOCHS : 9/10 Loss : 0.900666
The training accuracy is 60.99%
The validation loss is 0.4841
The valudation accuracy is 83.14%

EPOCHS : 9/10 Loss : 0.893229
The training accuracy is 61.41%
The validation loss is 0.4815
The valudation accuracy is 83.27%

EPOCHS : 10/10 Loss : 0.887517
The training accuracy is 61.80%
The validation loss is 0.4666
The valudation accuracy is 83.78%

EPOCHS : 10/10 Loss : 0.872435
The training accuracy is 62.16%
The validation loss is 0.4633
The valudation accuracy is 83.90%

```
[37]: optimizer = optim.Adam(AlexNet3.parameters(),lr=1e-6)
      train(AlexNet3,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",2,"d")
```

0%| | 0/3130 [00:00<?, ?it/s]

EPOCHS : 1/10 Loss : 1.000398
The training accuracy is 64.65%
The validation loss is 0.4515
The valudation accuracy is 84.18%

EPOCHS : 1/10 Loss : 0.840051
The training accuracy is 67.60%
The validation loss is 0.4455
The valudation accuracy is 84.52%

EPOCHS : 2/10 Loss : 0.841163
The training accuracy is 68.52%
The validation loss is 0.4427
The valudation accuracy is 84.39%

EPOCHS : 2/10 Loss : 0.834193
The training accuracy is 69.07%
The validation loss is 0.4453
The valudation accuracy is 84.51%

EPOCHS : 3/10 Loss : 0.842139
The training accuracy is 69.33%
The validation loss is 0.4452
The valudation accuracy is 84.52%

EPOCHS : 3/10 Loss : 0.832247
The training accuracy is 69.59%
The validation loss is 0.4508

The valudation accuracy is 84.45%

EPOCHS : 4/10 Loss : 0.829938

The training accuracy is 69.76%

The validation loss is 0.4495

The valudation accuracy is 84.32%

EPOCHS : 4/10 Loss : 0.840508

The training accuracy is 69.87%

The validation loss is 0.4373

The valudation accuracy is 84.74%

EPOCHS : 5/10 Loss : 0.830376

The training accuracy is 69.95%

The validation loss is 0.4497

The valudation accuracy is 84.30%

EPOCHS : 5/10 Loss : 0.828192

The training accuracy is 70.06%

The validation loss is 0.4463

The valudation accuracy is 84.54%

EPOCHS : 6/10 Loss : 0.817075

The training accuracy is 70.18%

The validation loss is 0.4493

The valudation accuracy is 84.43%

EPOCHS : 6/10 Loss : 0.840949

The training accuracy is 70.18%

The validation loss is 0.4396

The valudation accuracy is 84.62%

EPOCHS : 7/10 Loss : 0.830831

The training accuracy is 70.21%

The validation loss is 0.4396

The valudation accuracy is 84.62%

EPOCHS : 7/10 Loss : 0.815412

The training accuracy is 70.30%

The validation loss is 0.4503

The valudation accuracy is 84.29%

EPOCHS : 8/10 Loss : 0.811264

The training accuracy is 70.35%

The validation loss is 0.4431

The valudation accuracy is 84.75%

EPOCHS : 8/10 Loss : 0.815634

The training accuracy is 70.40%
The validation loss is 0.4378
The valudation accuracy is 84.83%

EPOCHS : 9/10 Loss : 0.819109
The training accuracy is 70.46%
The validation loss is 0.4370
The valudation accuracy is 84.89%

EPOCHS : 9/10 Loss : 0.816168
The training accuracy is 70.52%
The validation loss is 0.4373
The valudation accuracy is 84.64%

EPOCHS : 10/10 Loss : 0.823191
The training accuracy is 70.53%
The validation loss is 0.4414
The valudation accuracy is 84.70%

EPOCHS : 10/10 Loss : 0.817268
The training accuracy is 70.57%
The validation loss is 0.4332
The valudation accuracy is 84.92%

```
[38]: optimizer = optim.Adam(AlexNet3.parameters(),lr=1e-8)
      #AlexNet.load_state_dict(torch.load("Weights/part-d-run-2-Best.pth"))
      train(AlexNet3,train_dataloader,val_dataloader,10,20,loss_fn,optimizer,"cuda",3,"d")
```

0%| | 0/3130 [00:00<?, ?it/s]

EPOCHS : 1/10 Loss : 0.991773
The training accuracy is 64.99%
The validation loss is 0.4360
The valudation accuracy is 84.82%

EPOCHS : 1/10 Loss : 0.805074
The training accuracy is 68.39%
The validation loss is 0.4337
The valudation accuracy is 84.83%

EPOCHS : 2/10 Loss : 0.811057
The training accuracy is 69.44%
The validation loss is 0.4410
The valudation accuracy is 84.85%

EPOCHS : 2/10 Loss : 0.822433
The training accuracy is 69.90%
The validation loss is 0.4325

The valudation accuracy is 84.87%

EPOCHS : 3/10 Loss : 0.820671

The training accuracy is 70.17%

The validation loss is 0.4350

The valudation accuracy is 84.86%

EPOCHS : 3/10 Loss : 0.810344

The training accuracy is 70.41%

The validation loss is 0.4347

The valudation accuracy is 84.86%

EPOCHS : 4/10 Loss : 0.813434

The training accuracy is 70.58%

The validation loss is 0.4321

The valudation accuracy is 84.92%

EPOCHS : 4/10 Loss : 0.816192

The training accuracy is 70.63%

The validation loss is 0.4329

The valudation accuracy is 84.91%

EPOCHS : 5/10 Loss : 0.803453

The training accuracy is 70.76%

The validation loss is 0.4387

The valudation accuracy is 84.91%

EPOCHS : 5/10 Loss : 0.815467

The training accuracy is 70.82%

The validation loss is 0.4337

The valudation accuracy is 84.91%

EPOCHS : 6/10 Loss : 0.811431

The training accuracy is 70.85%

The validation loss is 0.4321

The valudation accuracy is 84.94%

EPOCHS : 6/10 Loss : 0.815065

The training accuracy is 70.88%

The validation loss is 0.4348

The valudation accuracy is 84.96%

EPOCHS : 7/10 Loss : 0.806256

The training accuracy is 70.94%

The validation loss is 0.4310

The valudation accuracy is 84.98%

EPOCHS : 7/10 Loss : 0.807493

The training accuracy is 70.99%
The validation loss is 0.4324
The valudation accuracy is 84.95%

EPOCHS : 8/10 Loss : 0.812556
The training accuracy is 71.01%
The validation loss is 0.4330
The valudation accuracy is 84.94%

EPOCHS : 8/10 Loss : 0.814946
The training accuracy is 71.05%
The validation loss is 0.4334
The valudation accuracy is 84.96%

EPOCHS : 9/10 Loss : 0.814383
The training accuracy is 71.06%
The validation loss is 0.4323
The valudation accuracy is 84.96%

EPOCHS : 9/10 Loss : 0.804518
The training accuracy is 71.10%
The validation loss is 0.4302
The valudation accuracy is 84.96%

EPOCHS : 10/10 Loss : 0.820406
The training accuracy is 71.08%
The validation loss is 0.4344
The valudation accuracy is 84.96%

EPOCHS : 10/10 Loss : 0.801487
The training accuracy is 71.12%
The validation loss is 0.4318
The valudation accuracy is 84.97%

[]: